


# Unity AI Gateway


**Your GenAI Control Plane –  
One place to access, govern, and  
observe all your AI traffic.**

---

August 2026



# Today, everyone is building with AI





There is  
**agent sprawl**

# Agent sprawl turns simple questions into governance issue across the C-suite



Head of AI

*How many agents are running, and who owns them?*



CISO

*Who can use an agent?*



Chief Data Officer

*What data did the agent access?*



Chief Risk Officer

*How can I keep an agent from deleting a data set?*



AI Platform Admin

*How should we monitor quality and debug issues?*



CFO

*What does each agent cost, and how much did we spend?*



# Why AI Gateway on Databricks

- 01 Open and Provider-Agnostic:** Route to any model from any provider with no lock-in. The gateway is open source.
- 02 Secure and Production-Ready:** Uses the same identity, networking, and compliance setup you already have. Handles 50T+ tokens per week in production today.
- 03 Data Stays in Your Environment:** Every request, response, and cost signal logs to Unity Catalog. Build dashboards, automate with Jobs, query with Genie.
- 04 Simplify Vendor Management:** Access leading models through Databricks without separate procurement, legal reviews, or capacity management per provider.



# Unity AI Gateway

The same governance you built for data now covers models, MCPs, and agents.

- **Access control:** Manage permissions for models, MCP servers, and agents using same data RBAC.
- **Discovery:** Browse and search all models, tools, and agents across your organization in one catalog.
- **Policy enforcement:** Define a guardrail or governance policy once and apply it to all AI assets.
- **Lineage:** End-to-end tracing (data to LLMs to agents) provides full visibility for debugging, optimization, and compliance.

The screenshot displays the Unity AI Gateway Beta dashboard. At the top, it says "AI Gateway Beta" with a "Send feedback" link. Below this, it states "Govern and access LLM endpoints. Learn more". The dashboard is divided into several sections:

- Coding agents:** Manage usage and access for coding agents. It lists Cursor, Codex, Gemini, and Other Integrations.
- Featured models:** Frontier models hosted by Databricks. It lists GPT-5.2, Claude Opus 4.6, Gemini 3.1 Pro, and Gemini 3 Pro.
- Fastest response:** Output tokens per second. It lists GPT OSS 20B (240.6), GPT OSS 120B (232.1), GPT-5 Nano (125.0), and Gemini 3 Flash (76.4).
- Highest Traffic:** Total input and output tokens in the last 7 days. It lists Claude Opus 4.5 (109.1M), Claude Opus 4.6 (18.2M), Claude Sonnet 4.5 (17.4M), and Claude Haiku 4.5 (8.2M).

Below these sections is a search bar and filters for "Owned by me", "Created by", and "Gateway features". A table lists various AI assets with columns for Name, Rate limits, Usage tracking, Inference table, Total tokens (7d), Created by, and Last modified.


| Name             | Rate limits | Usage tracking | Inference table | Total tokens (7d) | Created by  | Last modified         |
|------------------|-------------|----------------|-----------------|-------------------|-------------|-----------------------|
| production_endp  | ⌘           | ●              | ⌘               | 0                 | Ahmed Bilal | Feb 20, 2026, 12:21 F |
| prod_sonnet      | ⌘           | ●              | ⌘               | 0                 | Ahmed Bilal | Feb 20, 2026, 09:50   |
| databricks-gemi  | ⌘           | ●              | ⌘               | 94.1K             | Databricks  | Feb 18, 2026, 04:00 I |
| databricks-clauc | ⌘           | ●              | ⌘               | 17.4M             | Databricks  | Feb 18, 2026, 02:55 I |
| marketing_team   | ⌘           | ●              | ⌘               | 0                 | Ahmed Bilal | Feb 18, 2026, 01:52 F |
| teama            | ⌘           | ●              | ⌘               | 0                 | Ahmed Bilal | Feb 18, 2026, 09:56 . |
| databricks-clauc | ⌘           | ●              | ⌘               | 0                 | Databricks  | Feb 16, 2026, 10:00 F |
| databricks-clauc | ●           | ●              | ⌘               | 0                 | Databricks  | Feb 12, 2026, 11:40 F |



# Control access to models, MCPs, agents

Unity Catalog is the system of record for permissions across models, MCPs, and agents.


- **Models:** Control which LLMs your org can use. Grant or revoke access per model, per workspace.
- **Tools:** One place to manage credentials, permissions, and discovery for every tool your agents use.
- **Agents:** No agent gets more access than the person who triggered it. Identity flows end-to-end through UC.



# Enforce safety policies with Guardrails

Keep every AI interaction secure and risk-free across your AI surface.


- **Custom guardrails:** Define any business rule as a natural-language policy, attach it to any endpoint. No code deployment.
- **Flexible enforcement:** Apply on input, output, or both. Reject unsafe content or mask sensitive data.
- **Full audit trail:** Every guardrail decision, action, and outcome logged to inference tables automatically.



# One view across all AI activity

A single pane of glass for everything your AI does in production.


- **Every interaction logged:** Requests, responses, tokens, latency, errors, user identity, and guardrail outcomes, all in a UC-governed Delta table.
- **Query with tools you already use:** SQL, dashboards, jobs, alerts. Join AI logs with business data for richer insights.
- **Access governed by Unity Catalog:** Same permissions as any other table, with fine-grained control over who sees what.



# Govern the data that powers your AI

Trustworthy AI starts with governed, high-quality data.


- **Lineage from data to AI:** Trace how tables, volumes, and features flow into agents.
- **Data quality monitoring:** Catch anomalies in data assets before they reach agents.
- **Impact analysis:** Changing a table or swapping a model? See every downstream dependency first.



# Continuously evaluate your agents

Measure and improve agent accuracy in production with MLflow.


- **Scorers at every level:** Built-in judges, custom business rules, and code-based checks.
- **Production monitoring:** Scorers run on live traffic. Same judges from dev, scoring real requests.
- **Governed by Unity Catalog:** Traces, evaluations, and feedback all land in UC-governed Delta tables. No separate pipeline.



# Keep AI reliable when models fail

Keep agent quality consistent, even when models go down.

- **Fallback routing:** If the primary model degrades, traffic auto-routes to a backup. No code changes.
- **Load balancing:** Distribute traffic across providers for resilience and consistent response quality.



# Track and control AI spend

Full visibility into AI spend with built-in budget controls.

- **Token-level cost attribution:** Per-request, per-user, per-endpoint tracking in UC-governed inference tables.
- **[Q2 Roadmap] Budgets and spend limits:** Set thresholds per team, app, or endpoint. Alerts fire and access restricts at threshold.
- **[Q2 Roadmap] Coding agents included:** Claude Code, Cursor, Gemini CLI all route through the same gateway. One cost view for every AI tool your teams use.

Total Policies

5

Exceeded

2

5 users blocked

With Enforcement

2

blocking or throttling


Q Search budgets...

| NAME                   | WORKSPACES  | TAGS    | USERS & GROUPS                    | STATUS                                   | WHEN EXCEEDED                      |
|------------------------|-------------|---------|-----------------------------------|------------------------------------------|------------------------------------|
| default-monthly-budget | All         | —       | all-users<br>per-user - monthly   | <div><div></div></div> 4 of 11 exhausted | <div><div></div>Block Access</div> |
| power-user-budget      | All         | —       | power-users<br>per-user - monthly | <div><div></div></div> 1 of 4 exhausted  | <div><div></div>Block Access</div> |
| prod-design-budget     | All         | —       | product-design<br>shared - weekly | <div><div></div></div> 15% spent         | <div><div></div>Warn Only</div>    |
| dev-gateway-budget     | development | —       | —                                 | <div><div></div></div> 38%               | <div><div></div>Warn Only</div>    |
| chatbot-weekly-budget  | All         | chatbot | —                                 | <div><div></div></div> 52%               | <div><div></div>Warn Only</div>    |




SEE IT IN ACTION

# Demo



### Cost Breakdown

Weekly Cost (USD) by Endpoint



Top Endpoints by Cost (USD)

Top Models by Cost (USD)

Top Users by Cost (USD)



FUTURE

# Road Map

# What's Coming

## Guardrails

Content filtering, PII detection, and custom LLM judges to enforce safety policies on every request.

Q1

## Custom APIs

Extend gateway support to wide variety of APIs, including batch APIs, image/video output, and other APIs

Q1

## Spend Tracking

Enable detailed tracking and cost observability for LLMs, including external models, at granular levels such as users, projects, and requests.

Q1



# What's Coming

## Budget and Cost Controls

Q2

Set monthly spend thresholds per team/group/endpoint. Get alerts and automatically block access when limits are reached.

## MCP and Tools Gateway

Q2

Manage MCPs and tools in AI Gateway with centralized, fine-grained permissions for both external and DB-hosted MCPs, along with traffic and usage monitoring.

## Central AI Gateway

Q2

One account-level gateway in Unity Catalog to govern and observe all GenAI traffic across workspaces.




# What's Coming: Unity AI Gateway PuPr

Unity AI Gateway moves to Public Preview with UC integration | End of Q2

**Unity Catalog Integration:** LLMs, MCPs, and Agents are 1st-class UC securables. Govern once, apply across all workspaces.

**Fine-grained permissions and ABAC:** GRANT & DENY for AI assets using same governance primitives as data. Control access per user, group, or context attribute.

**Service Policies:** Unified policy enforcement at every AI interaction. Define any business rule as a function and attach it to an AI endpoint, MCP, or agent.



# What's Coming: Coding Agents Gateway Updates


Help admins govern and roll out coding agents across the org | Q2 2026

**Coding gateway rollout:** Roll out multi agent coding agent across your org with OAuth, MCP integration, and centralized config management.

**Per-user spend controls:** Set budgets and spend limits per developer across all coding tools in UC-governed usage tables.

**More OSS coding tools:** Hermes, OpenCode, Pi.

**Web search support for Claude:** Enable Claude code to use web search



# Unity AI Gateway Dashboards

| Option                                              | Claude model cost                                                                                                                          | Extra charge                                                                                                                     | What it means                                                                                                                                                               |
|-----------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                                     | Haiku 4.5: \$1 / \$5 per MToken; Sonnet 4.6: \$3 / \$15; Opus 4.8: \$5 / \$25.                                                             | None                                                                                                                             | Baseline reference.                                                                                                                                                         |
| Anthropic via Unity AI Gateway as an external model | Anthropic still charges the model invocation directly.                                                                                     | No extra Databricks compute/network fee; only AI Gateway feature billing applies, and current beta says those features are free. | <b>Today this is effectively price-parity with direct Anthropic.</b>                                                                                                        |
| Databricks-hosted Claude via Foundation Model APIs  | Priced by Databricks under Proprietary Foundation Model Serving, with Anthropic-specific pay-per-token / PT rates published by Databricks. | Separate from AI Gateway governance charges.                                                                                     | Internal DSV guidance was that Sonnet 4.6 was roughly at parity without in-geo, and about 10% higher only when in-geo was enabled on Azure Databricks discount assumptions. |



# Budgets (Preview) – Account Level

## Unity AI Gateway Budget

The screenshot shows the Databricks 'Usage' page with the 'Budgets' tab selected. The left sidebar contains navigation links: Workspaces, Catalog, Usage (highlighted with a red '1'), User management, Security, and Resilience. The main content area has tabs for Consumption, Cost Overview, Preview, Budgets (highlighted with a red '2'), and Consumption (Legacy). Below the tabs is a search bar and a table with columns: Name, Workspaces, Resource type, Tags, Shared threshold status, and Per user threshold status. A 'Create budget' button (highlighted with a red '3') is in the top right corner.

The 'Create budget' form is shown with the following sections:

- Scope:** Cannot be changed after creating budget.
- Name:** Budget
- Workspaces:** Select workspaces to include in this budget. All X
- Resource types:** Unity AI Gateway
- Resource tags (optional):** Select tags to restrict budgets to specific types of usage. Usage with all selected tag keys applied will be tracked. If a tag key has multiple values, for each key, usage with at least one of the selected values applied will be tracked. Select tags
- Shared thresholds:** Apply across all scoped spend.

| Users       | Monthly threshold | When exhausted |
|-------------|-------------------|----------------|
| (All users) | \$10.00           | Alert (1)      |

[+ Add threshold](#)
- Per-user thresholds:** Apply for each user, for their scoped spend. This threshold will be applied for all users.[+ Add threshold](#)

Buttons: Cancel, Create

**Databricks Budget Notification - test\_acon**

**Databricks** <noreply@databricks.com> to me

Wed, Jun 10, 1:33 PM (20)

Hello,

Your account's current Databricks spend for budget "test\_acon" triggered an alert. We're informing you so that you can investigate further and take appropriate action.

**Budget Name: test\_acon**  
**Current Spend: \$1.63**  
**Budgeted Amount: \$1.00**  
**Percent of Budget Spent: 163.26%**  
**Percent of Budget Period Elapsed: 33.33%**  
**Budget Filters:**  
{ "workspace\_id": { "operator": "IN", "values": [7474652653995122] } } (Product: Unity AI Gateway)

For more usage details, please visit your [Databricks account console](#).

If you no longer wish to receive these notifications, please modify the budget or contact your Databricks administrator. (Budgetid: 2c8af1fc-9f34-4c7c-b9e2-96339811d255)

- The Databricks Team



# Budgets (Preview) – Account Level

## Unity AI Gateway Budget for GENIE

FE-AWS-Sandbox

Usage

Consumption Cost Overview Preview **Budgets** Preview Consumption (Legacy)

Filters Q

Name Workspaces Resource type Tags Shared threshold status Per user threshold status

Create budget

Usage > Budgets >

Create budget

Scope

Cannot be changed after creating budget

Name

Budget

Workspaces

Select workspaces to include in this budget

All X

Resource types

Unity AI Gateway

Resource tags (optional)

Select tags to restrict budgets to specific types of usage. Tags are case-sensitive and must exactly match how they appear in your usage data.

Key

Usage with all the selected tag keys applied will be tracked

Values

For each key, usage with at least one of the selected tag values applied will be tracked

Custom tag contains

Select tags

Shared thresholds

Apply across all scoped spend

Users

(All users)

Monthly threshold

\$0.00

When exhausted

Block Alert (1)

Per-user thresholds

Apply for each user, for their scoped spend. This threshold will be applied for all users.

tag: databricks-product: genie

Usage > Budgets >

Create budget

Scope

Cannot be changed after creating budget

Name

Budget

Workspaces

Select workspaces to include in this budget

All X

Resource types

Unity AI Gateway

Resource tags (optional)

Select tags to restrict budgets to specific types of usage. Usage with all selected tag keys applied will be tracked. If a tag key has multiple values, for each key, usage with at least one of the selected values applied will be tracked.

Custom tag contains

databricks-product: genie

Select tags

Shared thresholds

Apply across all scoped spend

Users

(All users)

Monthly threshold

\$0.00

When exhausted

Block Alert (1)

Per-user thresholds

Apply for each user, for their scoped spend. This threshold will be applied for all users.

Cancel Create



Projects

# Migrations

# Migrating To New Gateway

1

## Nothing Breaks

Existing endpoints keep working. Run both in parallel and migrate at your own pace. Only change is updating the request URL.

2

## The Path Forward

The New AI Gateway is the primary place for all GenAI on Databricks. New use cases and endpoints should start here.

3

## Same Pricing

No billing changes with the New Gateway.



