

Databricks vs AWS native: executive capability map

This is an executive platform view, not a deep feature-by-feature teardown. The key pattern is that Databricks tends to package data, governance, SQL, AI and apps into one governed platform surface, while AWS usually covers the same outcomes through multiple strong point services that have to be connected together.

Simple readout

- Best Databricks displacement zones: governance, ETL/ingestion, AI/agents, and data/AI apps.
- AWS is strongest when a customer wants best-of-breed point services inside AWS, but the integration tax rises quickly as they try to unify data + AI + BI + security end to end.
- The AWS rows marked “Yes” below are the ones that typically require integration work, extra policy wiring, or manual setup/stitching.

Capability map

Capability	Databricks products	Closest AWS native tools	AWS needs manual setup / stitching?	Executive note
Unified governance for data and AI	Unity Catalog	Glue Data Catalog + Lake Formation + engine-specific layers across Athena, Redshift, EMR, SageMaker Lakehouse	Yes	This is one of the clearest Databricks advantages: AWS uses separate catalogs/security frameworks stitched together, while Databricks uses one catalog for data and AI assets.
Batch ingestion and ETL	Lakeflow Connect, Auto Loader, Auto CDC, Lakeflow	Glue ETL + EMR + MWAA/Airflow + Step Functions	Yes	AWS can do it, but the battlecards repeatedly call

Capability	Databricks products	Closest AWS native tools	AWS needs manual setup / stitching?	Executive note
	Jobs, Declarative Pipelines			out fragmented setup, separate orchestration, and manual incremental logic; Databricks positions this as one monitored flow.
Streaming and real-time pipelines	Structured Streaming + Lakeflow + Auto Loader/CDC	Kinesis + EMR/Glue streaming + orchestration	Yes	AWS has the pieces, but Databricks keeps processing, orchestration and observability on the same plane.
SQL warehouse and lake analytics	Databricks SQL + SQL Warehouses	Redshift + Athena	Partial	For pure SQL analytics AWS is viable, but once the customer wants one lakehouse across warehouse, ETL and AI, AWS starts layering multiple engines and catalogs.
BI and conversational analytics	AI/BI Dashboards, Genie Spaces, Genie One, business semantics	QuickSight for BI, typically paired with Amazon Q for NLQ	Usually	Databricks ties dashboards, semantics, governed access and natural-languag

Capability	Databricks products	Closest AWS native tools	AWS needs manual setup / stitching?	Executive note
				e exploration together on top of Unity Catalog; in AWS these concerns are usually split across BI, metadata and AI layers.
Traditional ML and MLOps	MLflow, Feature Store, Model Registry, Model Serving, AutoML	SageMaker AI + endpoints + experiments + surrounding AWS data services	Yes	Internal compete framing is that SageMaker's "unified" experience is still a wrapper around multiple underlying services, whereas Databricks keeps the ML lifecycle closer to the data platform.
GenAI and agents	Mosaic AI, AI Gateway, Vector Search, Model Serving, Agent Evaluation, Lakebase	Bedrock + AgentCore + Knowledge Bases / vector backends + IAM + optional SageMaker pieces	Yes	This is another major Databricks pitch: internal AI compete material describes AWS as a fragmented "kit of parts," especially around AgentCore, identity, memory,

Capability	Databricks products	Closest AWS native tools	AWS needs manual setup / stitching?	Executive note
				governance and evaluation.
Data and AI applications	Databricks Apps + Model Serving + Lakebase	Typically DIY on Lambda / API Gateway / App Runner / ECS / Cognito / Aurora or DynamoDB	Yes	Databricks Apps explicitly removes separate infrastructure for data/AI apps and inherits governance from the platform; this is where DIY AWS stacks often create extra platform work.
Federation / hybrid access	Lakehouse Federation and Catalog Federation	Athena federated access, Redshift Spectrum, Glue-backed hybrid catalog patterns	Partial	Databricks can read external DBs and catalogs without full migration, including AWS Glue catalog federation; useful when customers want coexistence instead of a big-bang move.
Data sharing and ecosystem exchange	Delta Sharing / OpenSharing / Marketplace	S3 bucket sharing, Lake Formation resource sharing, Redshift sharing	Partial to Yes	Databricks positions sharing as an open platform capability; in AWS, secure cross-account

Capability	Databricks products	Closest AWS native tools	AWS needs manual setup / stitching?	Executive note
				and cross-platform sharing often still falls back to bucket policies, network controls and IP allowlisting patterns.

Where the AWS integration tax is highest

If you want a quick talk track, the AWS tools most likely to require integration work or manual setup are:

1. Glue Data Catalog + Lake Formation when the customer wants one governance model across ETL, warehouse, BI and AI.
2. Glue + EMR + MWAA/Step Functions for ingestion and ETL orchestration.
3. Bedrock + AgentCore + memory/vector/governance layers for GenAI and agents.
4. DIY AWS app stacks when the app is tightly coupled to governed data and model access.
5. Cross-account / cross-platform sharing patterns that rely on bucket policies, egress IPs or additional network controls.

Executive positioning line

The cleanest way to frame this with a customer is: AWS gives you strong native components, but Databricks gives you the operating model that unifies them. The more the customer wants one governed platform for data engineering, warehouse, BI, ML, GenAI and apps, the more AWS turns into a stitching exercise and the more Databricks becomes the simplification story.

Sources

- [\[Platform Battlecard\] AWS Lakehouse](#)
- [\[Platform Battlecard\] AI on AWS](#)
- [Databricks Apps | Databricks on AWS](#)

- [Lakehouse reference architectures \(download\) | Databricks on AWS](#)
- [\[Platform Battlecard\] AWS Lakehouse](#)
- [Lakehouse reference architectures \(download\) | Databricks on AWS](#)
- [\[Use Case Battlecard\] Lakeflow Connect - Data Ingestion \(go/connect/battle\)](#)
- [05 AWS Competitive positioning](#)
- [Databricks AI/BI | Databricks on AWS](#)
- [Genie Spaces | Databricks on AWS](#)
- [\[Battlecard\] Databricks Apps](#)
- [Configure authorization in a Databricks app | Databricks on AWS](#)
- [Connect to external databases and catalogs | Databricks on AWS](#)
- [AWS Databricks](#)

Co-authored with Glean