

! README

Genie Spaces Deployment Guide

Databricks Genie Spaces

Enterprise Deployment Guide

What is this guide?

This guide is your roadmap for taking Databricks Genie Spaces from a successful proof of concept to a production-grade deployment that serves hundreds or thousands of business users across your organisation.

Whether you are just getting started with Genie Spaces or looking to scale an existing deployment, this guide provides practical, step-by-step guidance organised around six key pillars of enterprise readiness.

Note that Genie is the Databricks AI experience. Genie is part of a family of AI experiences built on the same governed data foundation:

- *Genie One: A simplified interface for business users to discover and interact with data assets.*
- *Genie Spaces: Domain-specific environments where data teams configure trusted data, metrics, and business rules that power Genie answers.*
- *Genie Code: The AI coding and data assistant for developers and technical practitioners inside the Databricks workspace.*

As of **May 20, 2026**, this guide is focused on Genie Spaces.

How to navigate this document

Section	What You Will Find	Time to Read
Fast Start	A high-level overview of the entire deployment approach — foundations, delivery, and scale-out	~5 minutes
Deployment Lifecycle	Six phases with key activities, exit criteria,	~5 minutes

Section	What You Will Find	Time to Read
	and suggested timelines	
The Success Pillars	The foundational areas to address when rolling out Genie Spaces at scale — even if some are already in place, revisit them to capture additional requirements	~5 minutes/pillar
Technical Checklist	A comprehensive checklist to validate before taking a Genie Space live	~5 minutes

Who should read this

Audience	What You Will Get
Business Analyst	Define the scope and requirements for Genie Spaces
Data & BI Engineer / Administrator	Set strong data foundations, build and maintain high-quality Genie Spaces Compliance, permissions, and access control requirements
Technical Executive Decision Maker	Strategy, governance, and rollout planning. Enable controlled enterprise AI adoption with governance
Business User	Understanding what Genie can do and how to get started

 Fast Start

Fast Start

Your 3-part deployment considerations

Enterprise Genie Space adoption rarely fails due to technology; it fails when organisations scale before their foundations are in place and teams are ready.


Success depends on aligning mindset and execution: focusing on clear business outcomes, demonstrating measurable ROI, and backing delivery with strong executive sponsorship, disciplined project management, and scalable foundations.

This section focuses on three essentials for a successful rollout: getting the foundations right, delivering spaces iteratively to drive early value, and sustaining quality as adoption scales.

Part 1: Start Solid with Foundation Checks

A [Genie Space](#) is a Databricks feature that allows business teams to interact with their data using natural language. It uses generative AI tailored to your organisation's terminology and data, with the ability to monitor and refine its performance through user feedback.

Every delivery plan includes an upfront preparation phase. The Success Pillars provide guidance on how to approach each foundational area.




Start these activities early and allow enough time for completion. Treat security sign-off, user onboarding, and identity provisioning as one-time foundational tasks, but don't let them delay space development; run them in parallel as separate workstreams.

Part 2: Fast delivery through empowered owners

We recommend iterating in waves, domain by domain, and empowering builders to deliver value regularly. This has a cultural implication: learning to accept good enough quickly and to build on feedback rather than perfect delivery.

Planning your enterprise Genie roll-out


Part 2: Fast delivery through empowered owners



Wave by wave, domain by domain.

Empowered builders and fast flow **deliver value fast.**

The cultural contract? Learning to accept good enough fast, and building on feedback over perfect, eventually.



Each space follows the same practitioner lifecycle. Each build should take **1-4 weeks from zero**, with [best practices](#) in place. The loop does not end at launch, as monitoring, feedback, and curation keep the space improving continuously.

One rule of thumb above all others: keep each space to a small number of tables. Focused spaces are more accurate than broad spaces with 20+ tables. When scope creeps, create a new space rather than expanding the existing one.

Part 3: The scale-out model

Once you have multiple spaces in production, individual attention is not enough. Three practices keep a growing portfolio healthy.

Practice	Why it matters
Push left	Fix data quality and metadata upstream in Unity Catalog; every improvement benefits every space that shares that data, not just the one you're working on.
Monitor at scale	Set up portfolio-level monitoring early. As adoption grows, catching degrading spaces before users notice becomes critical. Use system tables and benchmarks, not just manual spot-checks.
Community of practice	Space authors sharing patterns, a shared benchmark library, and a quarterly space health review cadence keep quality high without a central bottleneck.


Planning your enterprise Genie roll-out

Part 3: The scale out model

1

Push left: Improve data for all


Encourage owners to solve challenges once, for everyone, but pushing improvements upstream.



2

Monitoring at scale

As your adoption grows, monitoring and keeping quality in check becomes an important. Set up practices early.



3

Community of practice

Establish communities of practice. Share learnings and best practices. Encourage ownership mentalities and champion success



The next sections of the guide provide the details you need to lay the foundations and put a successful deployment lifecycle in motion.

Deployment Lifecycle


Deployment Lifecycle

The table below shows how a typical Genie Space deployment progresses from initial discovery through enterprise-wide rollout. Each phase builds on the previous one, though some phases run in parallel.

Phase	Name	Duration	Purpose	Key Activities	Exit Criteria
P0	Discovery & Assessment	2–3 weeks	Understand landscape, use cases, and readiness.	- Discovery workshops with business stakeholders - Data readiness assessment - Stakeholder identification - Success criteria definition - Topology	- Use cases prioritised - Data readiness assessed - Stakeholder RACI finalised - Success metrics agreed
P1	Build & Technical Validation	2–4 weeks	Build & technically validate Genie Spaces.	- Data preparation - Space configuration - Instruction curation - Benchmark testing	- All tables have metadata and descriptions - Benchmark accuracy $\geq 80\%$ for 1–2 weeks - Sample questions added - Knowledge Store configured
P2	Business User Validation	1–2 weeks	Validate with real business users.	- Onboard 5 power users - Collect structured feedback - Iterate on instructions	- Power users sign off - Evaluation scorecard meets thresholds - Critical feedback addressed
P3	Access & Consumption Setup	1–2 weeks (parallel to 2)	Configure access, identity, and channels.	- Set up permissions – Configure identity provisioning - Test consumption channels (UI, dashboards, apps, chat integrations)	- All target users are correctly entitled - Consumption channels tested end-to-end

Phase	Name	Duration	Purpose	Key Activities	Exit Criteria
P4	Monitoring & Onboarding Setup	Parallel	Set up monitoring, alerting, and user enablement.	- Deploy monitoring dashboards - Configure alerts - Publish internal user guide - Establish feedback processes	- Monitoring operational - Alerts configured - User documentation published
P5	Staged Rollout	4–8 weeks	Gradually scale from pilot to enterprise-wide rollout.	- Week 0: 10 users (validate stability) - Week 2–3: 50 users (department scale) - Week 4–5: 200 users (division scale) - Week 6+: enterprise-wide (full rollout)	- At each stage: usage monitored, feedback reviewed, instructions iterated - System performance is confirmed before expanding to the next stage

For a per-space operational tracker that maps to phases P1–P5 with milestones, owners, and exit criteria, use the Checklist subtab.



Technical Onboarding Checklist

Genie Space Technical Onboarding Checklist for <Customer>

This checklist supports the onboarding of a validated Genie Space into Production, ensuring smooth adoption and effective use of data insights. It outlines five key phases: validation by stakeholders, access setup, interface readiness, monitoring and onboarding, and a staged rollout plan. Follow each step, meet the exit criteria before advancing, and work closely with your Databricks team, especially as usage scales, to ensure a successful launch. This assumes you've already gone through Phase 0 - Discovery & Assessment.

Milestone Tracker				
✓ Milestone	👤 Owner	📅 Due Date	🔄 Status	📝 Notes / Blockers
Phase 1 - Technical User Validation of the Genie Space	👤 Person	📅 Date	Not started ▾	
Phase 2 - Business User Validation of the Genie Space	👤 Person	📅 Date	Not started ▾	
Phase 3 - Prepare Access and Consumption Layer	👤 Person	📅 Date	Not started ▾	
Phase 4 - Put Onboarding & Monitoring Processes in Place	👤 Person	📅 Date	Not started ▾	
Phase 5 - Stage the Rollout to More Users	👤 Person	📅 Date	Not started ▾	

Phase 1 - Technical User Validation of the Genie Space

Once you have a Genie Space that your creating team is happy with, we recommend you start the final validation process. For the purpose of this phase, we recommend that you actively use the Space yourselves as Space administrators, adjust the context as needed ([docs](#)), and regularly update and re-run benchmarks if any new questions arise last-minute before the next phase.

- ☐ A [limited](#) number of tables (the fewer the better)

- ☐ Tables should have Primary & Foreign Keys set, and table metadata is well [documented](#)¹ - columns have clear and precise names and descriptions²
- ☐ [Entity matching](#) (formerly “Value dictionaries”) helps when users’ prompts do not match the structure or values in the data (e.g., a user prompts “Florida” but the row values are “FL”)
- ☐ [Example SQL](#) is in place³ for the most common types of questions (preferably 10 or more). Parameterised queries are preferred because they can improve answer quality. Where a measure/filter/dimension can be expressed structurally, prefer a SQL Expression in the Knowledge Store over text instructions.
- ☐ [Instructions](#) mostly contain business terminology (acronyms, definitions, etc.), specific desired behaviours (e.g. if a prompt mentions X, always follow up until Y is also included), and miscellaneous guidance (e.g., always answer in language XYZ) that cannot be covered by table/column descriptions or Example Queries. If in doubt, always prioritise table/column descriptions and parameterised Example SQL.
- ☐ [Benchmarks](#) achieve the desired score: aim for ~80% accuracy; cover the top 15–20 expected user questions with 2–4 phrasings each (~30–80 benchmarks total); spaces support up to 500 benchmark questions; use the auto-suggested benchmarks feature to seed

Exit criteria:

- ☐ Benchmarks run at ~80% accuracy over the course of ~1-2 weeks
- ☐ Space Editors (technical stakeholders) have actively been using the space for ~1-2 weeks

¹ First and best place to put metadata is on table and column descriptions as only the relevant subset of those per prompt are sent, thus increasing the effective context available per prompt. Avoid describing column contents and how they should be used in General Text Instructions.

² Validate language consistency between table metadata and the preferred language of final users.

³ Use Example SQL to illustrate how different tables should be joined and filtered. The second best place for metadata is parameterized example SQL. Avoid describing calculations, filters and joins in General Text Instructions. Please note that you can also use comments in the Example SQL queries themselves to provide more context, for example, you can add alternate phrasing of your question that can map to the same SQL query.

Phase 2 – Business User Validation of the Genie Space

Since your benchmarks have been running regularly, it's time to onboard a handful (~5) of business power users (preferably those who were part of the initial scope definition as key stakeholders) for the business validation step. You don't need to set up the full access infrastructure just yet – you can give them temporary access to the workspace, with SQL entitlements, and the right permissions to use the space and run queries ([docs](#)).

- ☐ Business power users have had ~1-2 weeks of consistent testing and are satisfied with the answers at this stage of the Space
- ☐ Any changes needed to be made to the Space have been done while keeping in mind the steps outlined during [Phase 1](#) – for major changes, restart the timer on the ~1-2 weeks of testing

Exit criteria (can happen in parallel to Phases 3 and 4):

- ☐ Power users have tested the space for 1-2 weeks without any major changes needed to the space (you should anticipate ongoing review of user feedback and modification of space context for the lifetime of the space)

Phase 3 – Prepare Access and Consumption Layer

You can run some of the administration steps in this section in parallel to [Phase 2](#). You should aim to get in place a robust access management system in place, depending on the preferred way of sharing the Genie Space at scale within your organisation

- ☐ Choose the preferred way to consume the Genie Space – in the Genie UI, through a Dashboard Chat Box, or through Customer Web Apps or Teams/Slack integration via the Genie Conversation API⁴ (and build if required)
- ☐ Ensure basic [permissions](#) are in place for consuming users
 - ☐ Define credential methodology – viewer vs. embedded credentials – based on requirements. For the Genie Conversation API path, prefer OAuth over a single PAT token: OAuth preserves per-user UC privileges and per-user QPM accounting, whereas a shared PAT loses per-user visibility and is more likely to hit per-workspace limits.⁵

⁴ Please note that choosing a custom consumption layer such as a Web App or via Teams / Slack brings forth additional scope and complexity, which your technical team will have to manage on an ongoing basis, so ensure you keep that into consideration for your planning cycles.

⁵ Take into account considerations such as – do the users need to see personalised results (e.g. should only have access to a subset of the data in the Tables added to the Genie Space? If so, viewer credentials are required), do they need to see a simpler Databricks UI

- ☐ If required, set up [Automatic Identity Access](#) in Azure
- ☐ If required, set up [Share to Account Users \(embedded-credentials only\)](#) – still in Private Preview as of late March 2026, but being functionally superseded by the account-level Genie path below; prefer the GA path below for new builds
- ☐ For business-user consumption without workspace access, the preferred path is the account-level Genie experience (formerly “Databricks One”) – supports viewer credentials and a simplified UI on top of the Consumer entitlement; provision via AIM/SCIM. [The “Change default workspace access to Consumer access” toggle is in Public Preview \(April 2026\).](#)

Exit criteria (can happen in parallel to Phases 2 & 4):

- ☐ You have defined (and built, as needed⁶) the way in which business users will interact with Genie

Phase 4 - Put Onboarding & Monitoring Processes in Place

You can run some of the administration steps in this section in parallel to [Phase 2](#). Leverage tooling already available in Databricks to ensure that Space administrators can actively review the Space and react to issues in a timely manner, and make sure you publish this information and internal SLAs as part of onboarding users to the Genie Space:

- ☐ Set up SQL Warehouse monitoring to ensure SLAs are met. General SQL Warehouse monitoring guidance ([example blog](#))
- ☐ Set a schedule on which to address requests for [reviews](#)
- ☐ Set up monitoring of [Genie Audit events](#) (and alert [examples](#) available) to measure spikes in conversations/users - gauge adoption and whether you’re closing in on the current rate limits
- ☐ Leverage the Weekly Monitoring Digest in the Genie Space Monitoring tab, or consider this [community adoption dashboard](#) for Genie Space-specific observability
- ☐ Define the internal processes of how new users will get access to consume the Genie Space

that caters to business users, is the process to add users to a Databricks account / workspace currently complex, etc. Ensure you cover the differences between *Sharing to Account Users* and the *Consumer Role*.

⁶ If you’re looking for the Genie space to be consumed in a different Databricks workspace (e.g. moving from a DEV to a PROD workspace), as of January 2025, we have a [Management API](#) that can help with programmatic deployment between workspaces.

Exit criteria (can happen in parallel to Phases 2 & 3):

- ☐ An internal document available to the upcoming users of the Genie Space that provides them with details about the Genie Space and its scope, how to request access, [how to](#) use Genie to explore business data, as well as the cadence of feedback reviews
- ☐ Alerts are in place to notify the admin team when the SQL Warehouse is breaking defined SLAs
- ☐ Alerts are in place to notify the admin team when usage is reaching the current QPM limits (20 QPM for UI, 5 QPM for API)

Phase 5 - Stage the Rollout to More Users

Validate that Phases 1-4 have been completed before beginning the broader rollout to ensure end users have a positive experience. Define your desired end-state in terms of # of users that will have access to the Genie Space, as well as what the peaks of concurrency are that you think you'll reach (e.g. how does peak traffic look during your busiest hours), and work from those numbers backwards. For example, if you think the Genie Space will be accessible by 500 users in $T_0 + 2$ months, then create a staged roll-out plan that allows you to monitor via the solutions put in place in Phase 4, adjust the Space based on any unknown questions that come up via the new users, and keep up with the increasing number of manual reviews asked.

We strongly recommend that in the first few weeks, the admin team and the users onboarded work closely to solve any issues that were not preempted in the first two phases. We suggest the following staged approach (starting from **Apr 10, 2025** as T_0), and only progressing to the next step if there are no changes to the Space - otherwise, allow for a couple of days in between to mitigate any unpredictable quality changes that could arise due to changes:

- ☐ T_0 : Reach 10 users in the Space
- ☐ $T_0 + 2-3$ weeks: Reach 50 users in the Space
- ☐ $T_0 + 3-4$ weeks: Reach 200 users in the Space

For subsequent batches, we recommend continuing onboarding in a staged fashion, but do exponential increases after the first couple of stages.

When you are getting close to any of the monitoring limits defined in Phase 4:

- ☐ For SQL Warehouse issues, increase WH size/autoscaling as needed. Follow [these](#) guidelines for SQL Warehouse T-shirt sizing.

- ☐ For increases in # of manual reviews needed, ensure you're staffing the admin team appropriately
- ☐ If approaching QPM limits via API, options exist:
 - Work with the account team on Provisioned Throughput options
 - Wait for Pay-As-You-Go pricing



Success Pillars

Why Genie?

Genie lets your business users ask questions in plain language and get answers from your governed enterprise data, no SQL skills required. It bridges the gap between data teams and business users, enabling:

- **Faster time to insight.** Business users get answers in seconds, not days.
- **Consistent metrics.** Every answer draws from the same governed, curated data.
- **Lower support burden.** Routine "Can you pull this data?" requests are handled directly by Genie.
- **Built-in trust.** Users can see the SQL behind every answer, and critical queries can be certified for guaranteed accuracy.

Genie isn't a generic chatbot; it is a specialised tool tailored to your organisation's unique data landscape. **A production-ready Genie is a governed, trusted, supportable self-service analytics product that:**

- Answers real business questions against curated enterprise data.
- Respects your existing data governance and security model.
- Gives business users confidence in the answers.
- Scales across multiple teams and domains.
- Has clear ownership, testing, monitoring, and support.

The key shift – in a PoC, the question is "Can Genie answer some good questions?", whereas in production, the question is "Can Genie answer the right questions, consistently, securely, and at scale?".

The most common failure in enterprise Genie adoption is not technical: it is scaling before the foundations are in place and the people are ready. The most common success pattern is getting your foundations right, delivering spaces iteratively, and sustaining quality as you scale.

The Six Pillars

The six pillars cover everything an enterprise needs to deploy Genie successfully. Pillars 1 to 5 are foundational; Pillar 6 is optional and only applies if you need to extend Genie outside the native Databricks UI (Teams, Slack, custom apps, embedded experiences).

Pillar	Focus Area	What "good" looks like
1. Business Scope & Space Design	Design focused spaces, define a curation strategy, and assign ownership	Five or fewer tables per space; clear owner; 20 to 50 in-scope questions defined; Trusted Assets in place for critical KPIs
2. Data Foundation	Ensure data quality, metadata, and documentation are production-ready	Curated Silver / Gold tables in Unity Catalog; every column has a description; entity matching configured for categorical columns
3. Governance & Security	Configure access controls, credential models, and compliance requirements	Unity Catalog permissions, row filters, and column masks enforced per user; credential model documented per deployment pattern
4. Enablement & Change Management	Train authors and users, build trust, drive adoption	Space authors trained; sample questions and a prompt guide published; feedback loop active; common objections answered
5. Monitoring & Continuous Improvement	Benchmark accuracy, review feedback, iterate on quality	≥80% benchmark accuracy before production; Monitor tab reviewed on schedule; audit-log alerts configured
6. Integration & Embedding (Optional)	Extend Genie to Teams, Slack, APIs, and consumer portals	Throughput planned against observed demand; API integrations use minimum-privilege identity; channel choice matches the user audience

Through-lines across the pillars

A few themes run across all six pillars and are worth stating explicitly:

- **Foundations before features.** A focused, well-documented, governed space outperforms a broad space with clever instructions. Pillars 1 and 2 carry more weight than any prompt-engineering trick later in the stack.

- **Unity Catalog is the security boundary, not the Genie Space.** Row filters, column masks, and table grants automatically enforce per-user data access. A Genie Space defines the curated experience; it is not a hard data boundary.
- **Production-ready is measurable.** The $\geq 80\%$ benchmark-accuracy target is the gate to production. Below that, a space is not ready regardless of how it feels in a demo.
- **Field-recommended versus Databricks-published.** Most adoption targets and review cadences in this guide are field-recommended starting points. Only a high-accuracy gate and the documented throughput limits are published by Databricks. Calibrate the rest to your organisation.
- **Ownership beats one-off effort.** Spaces that drift in accuracy do so because no one owns them. A senior analyst or data product owner per space is the strongest predictor of sustained quality.

Top five ways enterprise rollouts stall

1. **Cross-domain spaces.** Trying to make one space answer everything. Always less accurate than three focused spaces.
2. **Skipping benchmarks.** Promoting to production without a benchmark suite leaves accuracy invisible and degradation undetectable.
3. **Treating Genie as a feature, not a product.** No owner, no backlog, no cadence equals quality decay within weeks of launch.
4. **Confusing the UI with the security boundary.** Hiding columns in the Knowledge Store is not access control. Unity Catalog grants are.
5. **Scaling channels before the space is production-quality.** Adding Teams or Slack to a marginal space amplifies the wrong answers to more users.

1. Business Scope and Space Design

Pillar 1: Business Scope and Space Design

Design each Genie Space for a defined business domain and a known set of user questions. Include the business terms, metrics, and tables required for that domain, and keep unrelated topics out of scope.

Scope limits answer quality. A space spanning unrelated domains is consistently less accurate than one that stays within its boundaries.

1.1 Design Principles

Keep the space narrow. Aim for [five or fewer tables per space, with 30 as the hard limit](#).

Use one space per business domain. Create separate spaces for Finance, Sales, Operations, and HR rather than one cross-domain space. Narrower spaces consistently produce better answers.

Before configuration starts, define:

- The audience.
- The 20 to 50 business questions the space should answer.
- The KPIs and datasets it owns.
- The topics explicitly out of scope.

State out-of-scope topics in the space instructions. Treat those instructions as guidance for Genie's behaviour, not as access control. Unity Catalog permissions still govern what users can actually query.

Size by audience:

- **Executive space:** 2 to 3 summary tables with approved metrics.
- **Domain analyst space:** 3 to 5 detailed tables for follow-up questions.

Validate the first 10 to 20 questions against the expected SQL and answers before opening the space to a wider audience. Expand from there.

Give the space a clear operating model:

- Owner
- Backlog

- Release cadence
- Support model
- Success metrics

1.2 Curation Hierarchy

A Genie Space improves through layered curation. Use the most structural layer that can express each rule. Reserve text instructions for behaviour that cannot be modelled in metadata or SQL.

- **Unity Catalog metadata:** The foundation. Accurate table and column descriptions, primary and foreign keys, and tags help Genie identify the right objects, and the same investment improves every other consumer of the catalog.
- **Metric Views:** Live in Unity Catalog, not in the space. They give a deterministic KPI definition: the calculation is authored once and applied wherever the view is referenced. Any synonyms declared in the metric view's agent metadata are available to Genie. Prefer Metric Views over space-level SQL expressions when the KPI is shared across multiple spaces.
- **SQL expressions:** Encode business logic that Genie can use during query generation. Define measures, filters, and dimensions as named expressions so the same logic is reused rather than reinvented per question.
- **Parameterised example SQL queries:** Map common question patterns to verified queries. When Genie matches a question to an example, it can use that query directly.
- **Text instructions:** Handle formatting, scope reminders, and clarification prompts. They iterate quickly but are open to interpretation, and the Databricks best-practices guidance recommends [using them as a last resort](#) when no structural option fits.

1.3 Knowledge Store Configuration

The [Knowledge Store](#) holds the space-specific configuration Genie uses to interpret tables, joins, calculations, and user phrasing in that space.

It contains:

- **Space-level metadata customisation:** Descriptions for tables, columns, business terms, and synonyms scoped to this space.

- **Space-level data customisation:** Simplified or focused views of the underlying data without changing the Unity Catalog tables.
- **Prompt matching:** Examples that resolve user phrasing to canonical entities, correct common spelling variants, and set output formats.
- **Join relationships:** Defined table relationships so JOIN logic is consistent.
- **SQL expressions:** Curated measures, filters, and dimensions.

Knowledge Store changes apply only to the current Genie Space. They do not modify Unity Catalog or any other Databricks asset.

1.4 Sample Questions

Configure 6 to 10 sample questions per space. They set user expectations on first open and provide working phrasing to start with.

Use the actual metrics, dimensions, and entities the space owns. Write literal questions, not templated placeholders. For a revenue space:

- What was revenue for Q1 2026?
- Who are the top 10 customers by bookings this year?
- How does gross margin compare across regions?
- Show monthly order volume for the last 12 months.
- Which regions missed the plan last quarter?
- Break down the pipeline by segment and geography.

Substitute the equivalents for your domain. Where possible, mirror the question patterns covered by your trusted assets so users land on trusted responses from their first interactions.

2. Data Foundation

Pillar 2: Data Foundation

Genie depends on clear metadata, consistent structures, and governed data to turn natural-language questions into accurate SQL queries. It uses your table and column metadata to identify the right dataset, measures, dates, and joins.

For a question such as "What were total sales last quarter?", Genie must understand:

- Which table contains the sales data?
- Which column holds the sales amount?
- Which column defines the date?
- How the relevant tables relate to each other.

If that context is missing or unclear, Genie is more likely to generate incorrect SQL. **Validate each space with representative benchmark questions before production.**

2.1 Data Layer Recommendations

Use cleaned, curated data rather than raw ingestion tables.

- **Raw / Bronze: avoid.** Quality issues, cryptic column names, and incomplete records make Genie unreliable.
- **Cleaned / Silver: use.** Validated, standardised, consistent data types.
- **Aggregated / Gold: prefer.** Optimised for analytics, clearer business logic, pre-calculated metrics.

2.2 Preparation Checklist

Preparation happens at two levels: each table in Unity Catalog, and then the Genie Space that uses them.

Per table (Unity Catalog)

Must have

- ☐ Register the table in Unity Catalog.
- ☐ Add a table comment that explains the contents and business context.
- ☐ Add a clear business description for every column.

- ☐ Use self-explanatory column names. Prefer `customer_lifetime_value` over `clv`.
- ☐ Store date fields as `DATE` or `TIMESTAMP`, not strings.

Should have

- ☐ Define primary and foreign key relationships.
- ☐ Apply Unity Catalog governance tags for PII and sensitivity classification.

Good to have

- ☐ Run `OPTIMIZE`, unless [predictive optimisation](#) is enabled.
- ☐ Document the data freshness cadence.

Per space (Knowledge Store and ACLs)

Must have

- ☐ Apply Unity Catalog permissions on the underlying data, and configure Genie Space ACLs on the space itself.

Should have

- ☐ Keep each space to five or fewer tables where possible. Treat 30 tables as a hard limit.
- ☐ Pre-join frequently joined tables into views, or define join relationships in the [Knowledge Store](#).
- ☐ Enable entity matching on categorical columns (regions, statuses, product names) so Genie maps user phrasings to canonical values.
- ☐ Enable format assistance on important columns to give Genie representative sample values.

Good to have

- ☐ Hide unnecessary or confusing columns in the Knowledge Store.

2.3 Metric Views

[Metric Views](#) are useful when one KPI definition needs to apply consistently across teams, dashboards, and Genie Spaces. Define a metric once, then let users query it by region, time period, or product line without rewriting the calculation each time.

When Genie answers from a Metric View, it uses the approved definition rather than a locally recreated formula. Adopt Metric Views when you have a shared set of standardised KPIs referenced across multiple teams or reports.

2.4 Data Quality Gates

Before adding a table to a production Genie Space, verify the following:

Gate	Check	Target
Completeness	Nulls in critical columns	Team-defined threshold met, for example, under 5%.
Freshness	Data refresh occurs on schedule	Documented cadence met.
Accuracy	Metrics match trusted reports	Benchmark queries match agreed outputs.
Schema Stability	No unplanned breaking changes	Change management process in place.
Performance	Query response times are acceptable	Tables optimised for expected workloads.

3. Governance & Security

Pillar 3: Governance & Security

Genie respects [Unity Catalog](#) automatically: your existing access controls, row-level security, and data-classification policies apply to Genie responses without re-configuration. This pillar covers the deployment patterns, controls, and review steps that preserve that posture in production.

3.1 How Genie Executes Queries: Credential Models

Genie generates read-only SQL from natural-language questions and runs it on a SQL warehouse. The identity under which the query runs depends on the deployment pattern.

Important. A Genie Space defines what users see in the UI: the tables, instructions, and examples available to them. It is **not a hard data boundary**. Unity Catalog permissions, row filters, and column masks define what each user can actually query. Rely on Unity Catalog grants for access enforcement, not on space composition alone.

Credential behaviour by deployment pattern:

Deployment Pattern	Who Runs the Query	End-User Requirements	Key Governance Implication
Standalone Genie Space	Compute credentials are embedded in the space by the author who configures the SQL warehouse.	Consumer access (or Databricks SQL entitlement) plus SELECT on the underlying Unity Catalog objects. End users do not need direct warehouse permissions.	Unity Catalog enforces what each user can see. Row filters and column masks apply based on the querying user's identity, except when using the embedded credential option within the embedding preview .
Dashboard Companion Space	Depends on the dashboard's sharing model. With shared data permission, the Genie Space uses embedded credentials (the publisher's). Otherwise, viewer credentials apply.	Access to the published dashboard.	Genie can answer questions from the full dataset query results, including fields or rows that are not visually displayed in the chart. Only include fields and rows in dashboard datasets that you intend viewers to reach through natural language.
Conversation API (embedded apps)	Authentication via service principal (OAuth M2M), user OAuth token (OAuth U2M), or PAT, depending on your integration design.	Depends on the token identity and its Unity Catalog grants.	Design your API integration to use the minimum-privilege identity appropriate for the use case, and document the access model

			explicitly.
--	--	--	-------------

3.2 Unity Catalog: Your Primary Security Layer

All Genie queries are governed by Unity Catalog. This means:

Control	How It Works
Table and view grants	A Genie Space can only query objects on which the user has SELECT access.
Row filters	Row filters defined in Unity Catalog filter the rows Genie returns, based on the querying user's identity (for example, region-based access).
Column masks	Column masks defined in Unity Catalog mask or redact sensitive columns (PII, compensation, and similar) based on the querying user's identity.
Audit logging	Genie interactions are captured in two places: the conversation history on the space (questions, answers, feedback), and Unity Catalog system tables for audit events, data-access lineage, and executed SQL.

Design principle: grant each user group access to the minimum data they need. Use curated views to simplify and constrain the queryable surface.

3.3 Access Model Design

Recommended access tiers:

Role	Genie Permission	Data Access	Typical Entitlement
Space owner or curator	CAN MANAGE	Full access to space tables	Databricks SQL access
Business user, validator, tester	CAN RUN	Governed by Unity Catalog grants, row filters, and column masks	Consumer access
Viewer	CAN VIEW	Read-only, no query execution	Consumer access

3.3.1 Consumer access and Genie

Two related but distinct concepts:

- **Consumer access** is a workspace entitlement for non-technical business users. It activates [Genie](#), dashboards, Databricks Apps, and related AI/BI experiences without requiring authoring entitlements (Workspace access or Databricks SQL access).
- **Genie** (previously Databricks One) is the simplified UI that consumer-access users land in. It focuses on consumption: no SQL editor, no notebooks, only dashboards and Genie spaces.

Key benefits

- No additional compute cost. You pay only for SQL warehouse usage.
- Simplified interface for business users.
- Provision users via your identity provider (Azure AD, Okta, and similar).

3.3.2 Onboarding business users to Genie

1. **Provision users at the account level.**
 - a. Add users through Automatic Identity Management.
 - b. Ensure users exist in the Databricks account before assigning the workspace.
2. **Assign users to the right workspace with consumer access.**
 - a. Add the users to the target workspace.
 - b. Grant consumer access to Genie, dashboards, and related shared assets without authoring permissions.
3. **Enable a custom URL (optional).**
 - a. Optionally enable a custom URL for the workspace.
 - b. Recommended to simplify the login experience across workspaces.

3.3.3 Access Configuration Checklist

- ☐ Unity Catalog permissions reviewed for all Genie Space tables.
- ☐ Row filters applied where needed (region, department, and similar).
- ☐ Column masks applied for sensitive data (PII, compensation, and similar).
- ☐ Space permissions set.
- ☐ Audit logging enabled and reviewed.
- ☐ Dashboard dataset scope reviewed for companion spaces.
- ☐ Credential model documented for each deployment pattern.
- ☐ Consumer access is configured for business users.

3.4 Compliance, Privacy, and Data Handling

Genie's compliance posture depends on your platform configuration, deployment choices, and organisational policies. This is an operational governance checklist, not a substitute for a formal compliance review. Use it to scope what your security and legal teams should validate. The items below are organised into three tiers: what is inherited automatically, what depends on your configuration, and what requires formal review.

Tier 1: Inherited from your Databricks platform controls

These items apply automatically based on your workspace and account configuration.

Item	Description
Unity Catalog governance	All data access is governed by Unity Catalog grants, row filters, and column masks.
Audit logging	Genie interactions are captured in audit logs.
Network controls	IP access lists, Private Link , and other network-level controls apply to Genie traffic.
Encryption	Data at rest and in transit is encrypted per your workspace configuration.
No model training on your data; zero retention by partners	Databricks does not use customer prompts, responses, or data to train foundation models (applies to all Databricks AI assistive features). Partner-powered AI features additionally use zero-data-retention endpoints from model partners.

Tier 2: Dependent on your deployment and configuration choices

These items vary based on how you configure and deploy Genie.

Item	What to validate
Credential model	Which identity runs the query depends on the deployment pattern (see Section 3.1). Document and validate for each space.
Dashboard companion data scope	Genie can access the full dataset query results, not just what is displayed visually. Review dashboard datasets before enabling the companion Genie.

Cross-Geo processing	AI assistive features use Databricks Geographic Regions (Geos) . Some regions may require cross-Geo processing to be enabled. Validate your workspace's Geo configuration.
Conversation history storage	Genie conversation history is stored (encrypted) in the Databricks control plane. Review retention and access policies. Customer-managed keys are supported for Genie spaces created after April 2025.
Data sent to models	Genie sends only the data needed to fulfil the request: prompts, relevant table metadata and values, errors, and input queries. Full table contents are not sent.
Partner-powered AI features	Partner-powered AI features must be enabled at the account or workspace level. Enablement defaults differ for workspaces with the compliance security profile (CSP) .

Tier 3: Requires formal security and compliance review

Validate these items with your security, legal, and compliance teams before production deployment.

Item	Guidance
Regulated data (HIPAA, PCI, FedRAMP, etc.)	AI assistive features can be used with regulated data when the required controls are configured, including the compliance security profile. Confirm your specific requirements with your Databricks account team and your internal compliance reviewers.
Data residency requirements	Confirm that your workspace Geo configuration meets your data residency requirements. Review cross-Geo processing settings.
AI usage policy	Confirm that your organisation's AI usage policies permit AI-assisted analytics.
Download and export controls	Decide whether result downloads should be restricted for Genie users. The setting lives under Workspace > Settings > Security > "SQL results download".

Retention and deletion	Review conversation history retention policies and deletion capabilities.
------------------------	---

Important. Do not summarise compliance posture in internal documentation without linking to the current official [Databricks AI trust and safety documentation](#). Compliance details evolve, so always reference the latest source.

3.5 Space Ownership: Who Maintains What

Role	Responsibility	Who Should Own This
Space owner	Overall accuracy, instruction curation, and ongoing refinement	Senior analyst or data product owner in your organisation
Data steward	Table quality, metadata accuracy, and data freshness	Data engineering team
Security reviewer	Permissions audit, row-level security verification	Security or governance team
Feedback reviewer	Monitor user questions, identify places to improve	Space owner or designated analyst

Important. Assign an ongoing space owner within your organisation. Spaces that are not actively monitored and maintained degrade in accuracy over time as business terminology, underlying data, and user needs evolve.

4. Enablement & Change Management

Pillar 4: Enablement & Change Management

Successful adoption depends on users understanding the tool, trusting it, and knowing how to use it effectively. This pillar outlines the enablement activities and change management considerations for a Genie deployment.

4.1 Stakeholder Engagement

Different stakeholders need different messages:

Stakeholder	What They Need to Know	Engagement Approach
Executive Sponsors	Business value, ROI, strategic priorities served	Executive demo using their organisation's own data
Data Platform Leaders	Architecture, governance, integration with existing stack	Demo & technical deep-dive on Unity Catalog integration
Space Authors / Curators	How to build, test, validate, and maintain spaces	2-day hands-on workshop with their actual data
Business Users	How to ask questions and interpret results. Verify results, provide feedback.	30-minute introduction plus guided first experience
IT / Security	Compliance, permissions, data privacy, audit trail	Security architecture review with compliance documentation

4.2 Recommended Training Program

Training should be tailored to the role each group plays in the deployment.

For Space Authors / Curators (technical)

Module	Duration	Content
--------	----------	---------

Genie Fundamentals	2 hours	How Genie works, architecture, capabilities and limitations
Data Preparation	2 hours	Metadata, comments, views, Metric Views
Space Configuration	3 hours	Knowledge Store, SQL expressions, example SQL, joins, synonyms
Trusted Responses	1 hour	When and how to create trusted queries and functions
Benchmarking & Iteration	2 hours	Building benchmark suites, interpreting results, iterating
Monitoring & Operations	1 hour	Monitoring tab, audit logs, feedback loops

For Business Users (non-technical)

Module	Duration	Content
Getting Started with Genie	30 min	What Genie is, what it can and cannot do, how to ask questions
Asking Effective Questions	30 min	Tips for getting better answers, using business terms, being specific
Understanding Responses	20 min	How to read results, what "Trusted" means, when to verify
Giving Feedback	10 min	How to use thumbs up / thumbs down, how feedback improves the space

The training content is available on the Databricks training catalog under [AI/BI for Data Analysts](#). For tailored sessions and workshops, reach out to your Databricks team.

4.3 Prompt Guide Template

Create a prompt guide for each space with the following sections.

Good question patterns

- "What was total revenue by region last quarter?"

- "Show me the top 10 customers by lifetime value."
- "How does this month's churn rate compare to last month?"

Tips for better answers

- Be specific about time periods, filters, and metrics.
- Use the business terms defined in your space, for example, "ARR" rather than "annual recurring revenue" if that is the configured synonym.
- Start simple, then add complexity.

Out-of-scope topics

- Questions about data not in this space.
- Questions that require real-time data if the space uses batch-refreshed tables.

4.4 Building Trust

Build trust through four practices:

- **Transparency.** Show users the SQL behind every answer, and encourage them to verify against known reports.
- **Trusted assets.** Use [Trusted Assets](#) for critical KPIs so users can identify validated responses by the "Trusted" badge in the UI.
- **Feedback loops.** Make it easy to report incorrect answers and show that user feedback drives improvements.
- **Gradual expansion.** Start with a small group of power users, then expand as confidence grows.

4.4.1 Overcoming Common Resistance

What Users Say	Why They Say It	How to Address It
"I don't trust AI-generated answers"	Concern about accuracy and reliability	Expose the transparency of the process by describing the thinking steps directly in the Genie Space UI. Demonstrate that users can see the exact SQL generated, for full transparency. Highlight Unity Catalog governance.
"This will replace my role"	Fear of job displacement	Explain that Genie reduces time spent writing routine SQL, while

		analysts remain responsible for interpreting results and handling strategic work.
"My existing dashboards work fine"	Change resistance, switching costs	Explain that Genie can sit alongside existing dashboards. Dashboard Companion lets users ask follow-up questions on charts they are already viewing.
"How do I know the data is right?"	Concern about data quality	Metric Views provide consistent metric definitions across tools. Trusted Assets produce validated results. All queries run against the same governed data.
"Seems hard to maintain"	Concern about ongoing effort	Show the monitoring tab and benchmark automation. Spaces improve continuously from user feedback. The curation effort decreases as the space matures.
"Can Genie Spaces access data outside Databricks?"	Concern about data scope	Genie Spaces can only query data registered in Unity Catalog (managed, external, foreign, and view objects). They cannot natively reach arbitrary SaaS apps or external databases unless that data is surfaced into Unity Catalog first. Lakehouse Federation lets you register external systems (Snowflake, SQL Server, PostgreSQL, BigQuery, and others) as foreign catalogs that Genie can then query directly without data movement, although we recommend creating views on top.

4.5 Measuring Adoption Success

Track these metrics through the rollout. The benchmark-accuracy targets are based on Databricks-published guidance for production-ready Genie Spaces; the rest are field-recommended starting points that you should calibrate to your organisation.

Metric	3-Month Target	6-Month Target
--------	----------------	----------------

Trained space authors	5 to 10	15 to 25
Active business users	50	200+
Genie Spaces in production	3 to 5	10 to 15
Weekly active users	20	100+
Benchmark accuracy	≥80%	≥85%
User satisfaction (positive feedback %)	>70%	>80%

5. Monitoring & Continuous Improvement

Pillar 5: Monitoring & Continuous Improvement

Maintain and improve Genie quality after launch. This pillar covers the monitoring, feedback, benchmarking, and review practices that keep a space accurate and useful over time.

5.1 Monitoring Tools

Databricks provides built-in monitoring at multiple levels through the [Monitor tab](#), [audit logs](#), and [system tables](#).

5.1.1 Monitor tab

For space owners, the Monitor tab helps you:

- View every question users have asked.
- See which questions received positive or negative feedback.
- Identify patterns, including which questions the space struggles with.
- Review requests where users have flagged responses that need attention.
- Track weekly digest signals: message volume, active users, and thumbs-up / thumbs-down rates.

Important. Thumbs up / thumbs down do not automatically change the space. Authors review feedback and decide what to act on.

5.1.2 Audit logs

Audit logs (captured in `system.access.audit` with `service_name = 'aibiGenie'`) help you:

- Track who accessed which spaces and when.
- Monitor conversation volumes and user activity trends.
- Set up alerts for unusual patterns such as spikes in negative feedback, approaching rate limits, or rising error rates. Alerts are configured as Databricks SQL Alerts on queries against the audit system tables; there is no Genie-native alert mechanism.

5.1.3 System tables

System tables expose Genie activity for analysis and cost attribution:

- `system.query.history`: all Genie-generated SQL queries, their duration, and success or failure rates. Filter for Genie traffic with `WHERE query_source.genie_space_id IS NOT NULL`, or for a single space with `WHERE query_source.genie_space_id = '<space-id>'`.
- `system.billing.usage`: cost of Genie usage via SQL warehouse consumption.
- `system.compute.warehouse_events`: compute scaling events triggered by Genie query load.
- `system.access.audit`: per-event audit records (also surfaced in 5.1.2 above).

5.2 The Feedback-to-Curation Loop

User feedback is the primary signal for keeping Genie quality high over time. Treat it as a repeatable operating loop rather than an occasional clean-up task.

1. **Capture** feedback from the Monitor tab: thumbs-up / thumbs-down signals, review requests, and recurring unanswered questions.
2. **Triage** the feedback by separating one-off issues from repeatable gaps, and prioritise questions tied to important business workflows.
3. **Curate** the space by updating instructions, adding [Trusted Assets](#), and expanding entity matching when users introduce new terminology.
4. **Validate** the change by rerunning the relevant benchmark questions and checking whether the issue is resolved without breaking existing coverage.
5. **Apply and observe**. Record the change in the curation log, then monitor the same question patterns to confirm the improvement holds.

5.3 Benchmarking: Measuring Accuracy Systematically

Benchmarks are a quality-assurance framework for Genie Spaces. Use a benchmark suite of representative business questions, each with multiple phrasings, paired with their expected correct SQL answers. Run the suite against the space and measure the percentage of Genie's predictions that are correct.

Databricks [recommends](#) aiming for 10 to 20 representative questions, with 2 to 4 phrasings each (roughly 30 to 80 benchmarks total). The product supports up to 500 benchmarks per space.

5.3.1 How to Use Benchmarks

- Create a benchmark suite covering the most important question categories.
- Run the suite after every significant change to instructions or data.
- Track accuracy over time; it should trend upward.
- Use benchmarks as a gate: do not promote a space to production unless it achieves $\geq 80\%$ accuracy.
- Use **Auto-suggested benchmarks** to bootstrap the suite. Genie can suggest test questions based on the space's tables, columns, and SQL examples.

5.3.2 Key Metrics to Track

The $\geq 80\%$ benchmark-accuracy target is Databricks-published guidance. The other metrics below are field-recommended signals; calibrate the targets to your organisation.

Metric	How to Measure	What to Look For
Benchmark accuracy	Benchmark feature	Stable or improving over time; $\geq 80\%$ before production
Thumbs-up rate	Monitor tab	Trending upward
Questions per user per week	Audit logs (system.access.audit)	Growing adoption
Error rate	Monitor tab	Declining over time
Trusted asset coverage on top questions	Space configuration	Increasing for high-frequency / high-cost questions
Time-to-insight	User survey	Measurable improvement

5.3.3 Benchmark Categories to Cover

- Simple aggregations ("What is total revenue?")
- Time-based queries ("How did sales compare month over month?")
- Filtered queries ("Show me only enterprise customers")
- Multi-table joins ("Which products did our top customers buy?")
- Top-N rankings ("Who are our top 10 accounts by revenue?")
- Complex business logic ("What is our customer churn rate?")

5.4 The Continuous Improvement Checklist

- ☐ Monitor tab reviewed on schedule.
- ☐ Audit-log alerts configured (via Databricks SQL Alerts on `system.access.audit`).
- ☐ Benchmark re-run schedule established.
- ☐ Feedback-to-curation loop documented and active.
- ☐ Curation change log maintained.
- ☐ Quarterly space health review scheduled.

5.4.1 Example Review Cycle

The cadence below is a field-recommended starting point. Calibrate it to the maturity of the space and the volume of user feedback.

Weekly

- Review the Monitor tab for new question patterns.
- Check for questions with negative feedback.
- Note questions Genie could not answer.

Every two weeks

- Add new [Trusted Assets](#) for frequently asked questions.
- Update entity matching if users introduce new terminology.
- Address any review requests from users.

Monthly

- Run the full benchmark suite.
- Review accuracy trends: Is the space improving?
- Review trusted asset coverage on the top questions.
- Discuss feedback themes with business stakeholders.

Quarterly

- Conduct a full space audit; remove outdated instructions and update metric definitions.
- Review data freshness and schema changes.
- Reassess whether the scope of the space still matches user needs.
- Consider whether new spaces are needed for emerging use cases.

6. Integration & Enablement

Pillar 6 (Optional): Integration & Embedding

Genie Spaces work best when users can access them directly in Genie (previously Databricks One), as that provides the richest native experience. If your channel strategy requires users to stay in other tools, you can also extend Genie into dashboards, APIs, collaboration tools, and custom applications.

6.1 Consumption Channels

Channel	Best For	Setup Effort	Description
Genie Space UI (workspace)	Data teams, analysts	Immediate	Native Databricks experience with the broadest feature coverage
Dashboard Companion	Dashboard consumers	Immediate	Ask follow-up questions on published AI/BI dashboards through the companion Genie Space
Genie (consumer access)	Broad business rollout	Days	Simplified landing experience for business users, with no SQL knowledge required
Conversation API	Custom apps, portals, chatbots	2 to 3 weeks	Programmatic access for branded or embedded experiences
Microsoft Teams (Copilot Studio)	Organisations using M365	1 to 2 weeks	Low-code integration through the Databricks connector for Microsoft Power Platform
Microsoft Teams / Slack (custom bot)	Full UI control	3 to 4 weeks	Custom bot with branded workflows, prompts, and controls

Note. Databricks publishes workspace-level throughput limits for Genie, but not a fixed user-count threshold for external integrations. As field guidance, if you expect around 500 concurrent users or sustained high-volume API traffic, validate routing, SQL warehouse sizing, request queuing, and end-to-end latency early.

6.2 Choosing Your Integration Path

For most enterprise rollouts, we recommend a phased approach:

- Start with Genie UI and Dashboard Companion for initial validation. This gives you the richest feature set for testing and tuning.
- Add Genie consumer access for broad business rollout once the space is production-quality. This is usually the lowest-friction path to bring non-technical users onto Genie.
- Add Teams, Slack, or custom integrations after the space is production-quality, for specific business cases that need them.

6.3 Conversation API for Custom Experiences

The [Conversation API](#) lets you embed Genie in custom applications, chatbots, and portals. The same surface area also exposes the Genie Spaces management APIs for creating, configuring, backing up, and promoting spaces across workspaces.

Key capabilities

- Start and continue multi-turn conversations programmatically.
- Send questions and receive structured responses.
- Retrieve generated SQL, query results, message status, and metadata.
- Manage Genie Space configuration through the management APIs.

Authentication options

- [OAuth M2M for service principals](#), for server-side applications and unattended workloads.
- [OAuth U2M for users](#), when a browser-based user flow is available.

Best practices for API integration

- Poll message status every 5 to 10 seconds until you receive a conclusive state, such as COMPLETED, FAILED, or CANCELLED.
- Use exponential backoff for retries and polling, increasing the delay up to one minute for long-running queries.
- Start a new conversation for each user session to avoid unintended context reuse, which can reduce accuracy.
- Implement request queuing, retry logic, and logging in the calling application.
- Stop polling after about 10 minutes for most queries, then return a timeout or manual follow-up path.

- Only POST requests count towards the questions-per-minute limit; polling GET requests do not.
- Render charts in your own UI if needed. The API returns structured tabular results and SQL, not rendered visualisations.

Planning considerations

- Use the minimum-privilege identity model for each integration.
- Curate and benchmark the Genie Space before exposing it through the API.
- Plan SQL warehouse capacity separately from API rate limits. Every Genie query still runs on the attached warehouse.
- Expect some lag between new UI features and API availability.
- Use the management APIs to export and recreate space configuration across environments.

6.4 Environment Promotion

For enterprise deployments, keep promotion explicit and repeatable:

- Use the Genie Spaces management APIs to move space configuration between environments. If you already standardise delivery with [Declarative Automation Bundles](#) (previously Databricks Asset Bundles), keep the surrounding deployment workflow in the same promotion pipeline. Native Genie Space resources in DABs are still in development; use the Management API as the primary CI/CD path for spaces themselves.
- Maintain separate catalogues for dev, staging, and production data.
- Re-run benchmarks before every promotion.

6.5 Supervisor Agent

For large organisations with many Genie Spaces, the [Agent Bricks Supervisor Agent](#) can act as an alternate conversational entry point alongside Genie. Either option routes finance questions to a finance space, HR questions to an HR space, and so on. It can also coordinate Genie Spaces alongside Knowledge Assistant agents, Unity Catalog functions, MCP servers, and custom agents in one orchestrated experience.

This reduces navigation friction once an organisation has multiple domain-specific spaces. It also creates a cleaner front door for downstream applications, while preserving Unity Catalog access controls across the underlying spaces and subagents.

6.6 Throughput Planning

For API-based integrations, plan on observed demand rather than theoretical peak usage.

Key insight. In practice, the concurrent query rate is much lower than the peak user count.

This is because:

- Users pause to read and interpret results between questions.
- Load tends to arrive in short bursts rather than as a continuous stream.
- Application-side caching and queuing can reduce the number of repeated calls.
- Time-zone distribution often smooths demand in larger deployments.

Recommended approach

- Start with a small pilot cohort and measure actual query patterns.
- Monitor peak concurrent queries per minute over at least 30 days.
- Right-size capacity using observed traffic, not headline user counts.
- Scale incrementally as adoption expands.

Throughput tiers May 20, 2026

- Databricks UI: up to 20 questions per minute per workspace, across all Genie Spaces.
- Conversation API free tier: best-effort five questions per minute per workspace, across all Genie Spaces.
- Copilot Studio integrations inherit the Conversation API limit when calling Genie via the Databricks connector.
- For higher throughput, contact your Databricks account team. Dedicated capacity is available in increments above the free tier.