

<INSERT CUSTOMER LOGO>



Please leverage the more
comprehensive

 [CUSTOMER-READY] Genie Enterpri...

Delete the red text after reading (there's some at the end too). Replace yellow/blue highlights as needed.

This document is intended to be minimally adjusted and to serve as a customer-facing template when exposing a Genie Space to business users. Please note that, outside of this document, you should have already demonstrated **business value** in Genie for your customer and followed our [documentation](#) to kick-start the evaluation process, as that provides you with information on how to set up your Genie Space.

We want to move from “as a data team, we validated the room and want to roll out to the business” to “as a data team, we have collected sample questions from the business, we tuned the space and used it ourselves as new questions came trickling in, we rolled it out to select members and have iterated a few times, and only then we roll it out to the business.”. This will ensure the data team’s success and a favourable adoption of the Genie Space.

You can find the accompanying deck here

 [CUSTOMER-READY] Genie Space Onboarding Plan for <CUSTOMER>

If you’re looking to make suggestions to this document, please leave comments on

 [STAGING] Genie Space Rollout Plan Template

Replace Yellow Highlighted Items

Check Blue Highlighted Items for Updates (mainly previews)

Genie Space Onboarding Plan for <Customer>

This checklist supports the onboarding of a validated Genie Space into Production, ensuring smooth adoption and effective use of data insights. It outlines five key phases: validation by stakeholders, access setup, interface readiness, monitoring and onboarding, and a staged rollout plan. Follow each step, meet the exit criteria before advancing, and work closely with your Databricks team, especially as usage scales, to ensure a successful launch.

Milestone Tracker				
✓ Milestone	👤 Owner	📅 Due Date	🔄 Status	📝 Notes / Blockers
Phase 1 - Technical User Validation of the Genie Space	👤 Person	📅 Date	Not started ▾	
Phase 2 - Business User Validation of the Genie Space	👤 Person	📅 Date	Not started ▾	
Phase 3 - Prepare Access and Consumption Layer	👤 Person	📅 Date	Not started ▾	
Phase 4 - Put Onboarding & Monitoring Processes in Place	👤 Person	📅 Date	Not started ▾	
Phase 5 - Stage the Rollout to More Users	👤 Person	📅 Date	Not started ▾	

Phase 1 - Technical User Validation of the Genie Space

Once you have a Genie Space that your creating team is happy with, we recommend you start the final validation process. For the purpose of this phase, we recommend that you actively use the Space yourselves as Space administrators, adjust the context as needed

([docs](#)), and regularly update and re-run benchmarks if any new questions arise last-minute before the next phase.

- ☐ A [limited](#) number of tables (the fewer the better)
- ☐ Tables should have Primary & Foreign Keys set, and table metadata is well [documented](#)¹ - columns have clear and precise names and descriptions²
- ☐ [Entity matching](#) (formerly “Value dictionaries”) helps when users’ prompts do not match the structure or values in the data (e.g., a user prompts “Florida” but the row values are “FL”)
- ☐ [Example SQL](#) is in place³ for the most common types of questions (preferably 10 or more). Parameterised queries are preferred because they can improve answer quality. Where a measure/filter/dimension can be expressed structurally, prefer a SQL Expression in the Knowledge Store over text instructions.
- ☐ [Instructions](#) mostly contain business terminology (acronyms, definitions, etc.), specific desired behaviours (e.g. if a prompt mentions X, always follow up until Y is also included), and miscellaneous guidance (e.g., always answer in language XYZ) that cannot be covered by table/column descriptions or Example Queries. If in doubt, always prioritise table/column descriptions and parameterised Example SQL.
- ☐ [Benchmarks](#) achieve the desired score: aim for ~80% accuracy; cover the top 15–20 expected user questions with 2–4 phrasings each (~30–80 benchmarks total); spaces support up to 500 benchmark questions; use the auto-suggested benchmarks feature to seed

Exit criteria:

- ☐ Benchmarks run at ~80% accuracy over the course of ~1-2 weeks

¹ First and best place to put metadata is on table and column descriptions as only the relevant subset of those per prompt are sent, thus increasing the effective context available per prompt. Avoid describing column contents and how they should be used in General Text Instructions.

² Validate language consistency between table metadata and the preferred language of final users.

³ Use Example SQL to illustrate how different tables should be joined and filtered. The second best place for metadata is parameterized example SQL. Avoid describing calculations, filters and joins in General Text Instructions. Please note that you can also use comments in the Example SQL queries themselves to provide more context, for example, you can add alternate phrasing of your question that can map to the same SQL query.

- ☐ Space Editors (technical stakeholders) have actively been using the space for ~1-2 weeks

Phase 2 - Business User Validation of the Genie Space

Since your benchmarks have been running regularly, it's time to onboard a handful (~5) of business power users (preferably those who were part of the initial scope definition as key stakeholders) for the business validation step. You don't need to set up the full access infrastructure just yet - you can give them temporary access to the workspace, with SQL entitlements, and the right permissions to use the space and run queries ([docs](#)).

- ☐ Business power users have had ~1-2 weeks of consistent testing and are satisfied with the answers at this stage of the Space
- ☐ Any changes needed to be made to the Space have been done while keeping in mind the steps outlined during [Phase 1](#) - for major changes, restart the timer on the ~1-2 weeks of testing

Exit criteria (can happen in parallel to Phases 3 and 4):

- ☐ Power users have tested the space for 1-2 weeks without any major changes needed to the space (you should anticipate ongoing review of user feedback and modification of space context for the lifetime of the space)

Phase 3 - Prepare Access and Consumption Layer

You can run some of the administration steps in this section in parallel to [Phase 2](#). You should aim to get in place a robust access management system in place, depending on the preferred way of sharing the Genie Space at scale within your organisation

- ☐ Choose the preferred way to consume the Genie Space - in the Genie UI, through a Dashboard Chat Box, or through Customer Web Apps or Teams/Slack integration via the Genie Conversation API⁴ (and build if required)
- ☐ Ensure basic [permissions](#) are in place for consuming users
 - ☐ Define credential methodology - viewer vs. embedded credentials - based on requirements. For the Genie Conversation API path, prefer OAuth over a single PAT token: OAuth preserves per-user UC privileges and per-user QPM

⁴ Please note that choosing a custom consumption layer such as a Web App or via Teams / Slack brings forth additional scope and complexity, which your technical team will have to manage on an ongoing basis, so ensure you keep that into consideration for your planning cycles.

accounting, whereas a shared PAT loses per-user visibility and is more likely to hit per-workspace limits.⁵

- ☐ If required, set up [Automatic Identity Access](#) in Azure
- ☐ If required, set up [Share to Account Users \(embedded-credentials only\)](#) – still in Private Preview as of late March 2026, but being functionally superseded by the account-level Genie path below; prefer the GA path below for new builds
- ☐ For business-user consumption without workspace access, the preferred path is the account-level Genie experience (formerly “Databricks One”) – supports viewer credentials and a simplified UI on top of the Consumer entitlement; provision via AIM/SCIM. [The “Change default workspace access to Consumer access” toggle is in Public Preview \(April 2026\).](#)

Exit criteria (can happen in parallel to Phases 2 & 4):

- ☐ You have defined (and built, as needed⁶) the way in which business users will interact with Genie

Phase 4 - Put Onboarding & Monitoring Processes in Place

You can run some of the administration steps in this section in parallel to [Phase 2](#).

Leverage tooling already available in Databricks to ensure that Space administrators can actively review the Space and react to issues in a timely manner, and make sure you publish this information and internal SLAs as part of onboarding users to the Genie Space:

- ☐ Set up SQL Warehouse monitoring to ensure SLAs are met. General SQL Warehouse monitoring guidance ([example blog](#))
- ☐ Set a schedule on which to address requests for [reviews](#)
- ☐ Set up monitoring of [Genie Audit events](#) (and alert [examples](#) available) to measure spikes in conversations/users – gauge adoption and whether you’re closing in on the current rate limits

⁵ Take into account considerations such as – do the users need to see personalised results (e.g. should only have access to a subset of the data in the Tables added to the Genie Space? If so, viewer credentials are required), do they need to see a simpler Databricks UI that caters to business users, is the process to add users to a Databricks account / workspace currently complex, etc. Ensure you cover the differences between *Sharing to Account Users* and the *Consumer Role*.

⁶ If you’re looking for the Genie space to be consumed in a different Dabtricks workspace (e.g. moving from a DEV to a PROD workspace), as of January 2025, we have a [Management API](#) that can help with programmatic deployment between workspaces.

- ☐ Leverage the Weekly Monitoring Digest in the Genie Space Monitoring tab, or consider this [community adoption dashboard](#) for Genie Space-specific observability
- ☐ Define the internal processes of how new users will get access to consume the Genie Space

Exit criteria (can happen in parallel to Phases 2 & 3):

- ☐ An internal document available to the upcoming users of the Genie Space that provides them with details about the Genie Space and its scope, how to request access, [how to](#) use Genie to explore business data, as well as the cadence of feedback reviews
- ☐ Alerts are in place to notify the admin team when the SQL Warehouse is breaking defined SLAs
- ☐ Alerts are in place to notify the admin team when usage is reaching the current QPM limits (20 QPM for UI, 5 QPM for API)

Phase 5 - Stage the Rollout to More Users

Validate that Phases 1-4 have been completed before beginning the broader rollout to ensure end users have a positive experience. Define your desired end-state in terms of # of users that will have access to the Genie Space, as well as what the peaks of concurrency are that you think you'll reach (e.g. how does peak traffic look during your busiest hours), and work from those numbers backwards. For example, if you think the Genie Space will be accessible by 500 users in $T_0 + 2$ months, then create a staged roll-out plan that allows you to monitor via the solutions put in place in Phase 4, adjust the Space based on any unknown questions that come up via the new users, and keep up with the increasing number of manual reviews asked.

We strongly recommend that in the first few weeks, the admin team and the users onboarded work closely to solve any issues that were not preempted in the first two phases. We suggest the following staged approach (starting from **Apr 10, 2025** as T_0), and only progressing to the next step if there are no changes to the Space - otherwise, allow for a couple of days in between to mitigate any unpredictable quality changes that could arise due to changes:

- ☐ T_0 : Reach 10 users in the Space
- ☐ $T_0 + 2-3$ weeks: Reach 50 users in the Space
- ☐ $T_0 + 3-4$ weeks: Reach 200 users in the Space

For subsequent batches, we recommend continuing onboarding in a staged fashion, but do exponential increases after the first couple of stages.

When you are getting close to any of the monitoring limits defined in Phase 4:

- ☐ For SQL Warehouse issues, increase WH size/autoscaling as needed. Follow [these](#) guidelines for SQL Warehouse T-shirt sizing.
- ☐ For increases in # of manual reviews needed, ensure you're staffing the admin team appropriately
- ☐ If approaching QPM limits via API, options exist:
 - Work with the account team on Provisioned Throughput options
 - Wait for Pay-As-You-Go pricing

Please make sure you review

 [\[INTERNAL ONLY\] FAQs for AI/BI - Dashboards & Genie - go/aibi/faq](#)