

Databricks AI Assistive Features – Security FAQ

Last updated: Feb 11, 2026

Andrew Weaver Weston Hutchins Ted Tomlinson Samantha Banchik Suchismita Pahi

[go/databricksiq/security/faq](https://go.databricksiq/security/faq)

[go/databricksassistant/security/faq](https://go.databricksassistant/security/faq)

Databricks Internal Only (Review each question for instructions on sharing)

[Databricks AI assistive features – trust and safety is your best customer-facing resource!](#)

How to use this document

This document helps account teams who receive customer questions about DatabricksIQ powered features.

For a general understanding of platform security please see go/platform/security.

For answers to more general questions about the DatabricksIQ, please refer to the [Databricks AI assistive features – trust and safety, public documentation](#), and related FAQs ([\[Deprecated\] DatabricksIQ FAQ \(go/databricksiq-faq\)](#) , [Genie Code FAQ](#)).

 Databricks AI Features Security & Architecture is a customer-facing security deck.

Important: Do not share this entire document directly with your customers. Copy & paste specific answers, or use it to inform yourself before speaking with your customers.

Question classification: Some questions can be shared freely, but others are intended as reactive questions only if a customer asks for more detail.

New questions: Ask in the #security, #ai-bi-genie or #eng-genie-code-team slack channels

How to use this document

General Questions

- Q: What are Databricks AI-powered features?
- Q: What is Genie Code?
- Q: How do I enable Genie Code?
- Q: Is Genie Code enabled by default for workspaces?
- Q: What happens when I enable the Compliance Security Profile on my workspace? Will Genie Code be disabled automatically?
- Q: Are there plans for more fine-grained access controls (per-workspace, per-user), or is the feature only available per account?
- Q: If someone enables Databricks AI-powered features for my account or workspace is that captured in my audit logs?
- Q: Are the requests sent by a user to Genie Code captured in audit logs?
- Q: Are the questions sent by a user within an AI/BI Genie space captured in audit logs?
- Q: What is Genie Code Autocomplete?

Model Questions

- Q: What models are AI-powered features using?
- Q: Are there plans to integrate with other models?
- Q: Why doesn't Databricks support customer-provided models, like using a customer-provided OpenAI endpoint?
- Q: What is the difference between OpenAI and the Azure OpenAI Service?
- Q: If you change the underlying models or services between providers will I be notified?
- Q: Do you use the same Azure OpenAI subscription/service for all customers? How are my requests isolated from other customers?
- Q: Do you use the same Anthropic service for all customers? How are my requests isolated from other customers?
- Q: Do you use the same Databricks-managed model serving endpoints for all customers? How are my requests isolated from other customers?
- Q: How accurate is Genie Code?
- Q: Will Genie Code execute dangerous code?

Data Questions

- Q: What data is being sent to the models?
- Q: Does the metadata sent to the models respect the user's Unity Catalog permissions?

Q: Where is the data being stored?

Q: Is my data being used to train models?

Q: If I execute a query with results, and then ask a question, do the results of my query get sent to the model?

Q: Can I opt out of Azure OpenAI and Anthropic on Bedrock storing my data?

Q: Can I request an exemption from abuse monitoring and human review on Azure OpenAI and/or Anthropic on Bedrock?

Q: Is any content filtering applied to the responses to filter for harmful responses?

Q: Are there any guardrails in place to strip out PII/sensitive data before sending it to Databricks-managed models or our model partners?

Q: What telemetry is captured?

Q: Who owns the output data? If AI-powered features are generating some code, who owns that IP?

Q: Will AI-powered features generate copyrighted code? Am I legally protected if I use that in my production software?

Legal & Compliance Questions

Q: Are there any data residency considerations I should be aware of?

Q: Can I request for my data to be deleted?

Q: Can I use Genie Code with tables that contain regulated data (FEDRAMP/PHI/PCI/IRAP)?

Q: Can I enable/use Genie Code in Compliance security profile enabled workspaces?

Q: For AWS GovCloud and Azure Government customers, do we use Azure OpenAI and/or Anthropic instances in Azure Government?

Q: What legal terms govern the use of AI-powered features? Is it the Azure OpenAI API terms or other terms determined by Databricks?

Q: Does Databricks SOC2 report cover AI-powered features?

Q: Has there been a penetration test for AI-powered features?

Q: As part of the Microsoft Responsible AI review, what harm categories were assessed?

Encryption Questions

Q: Is everything encrypted in transit? Does any traffic hit the public internet?

Q: Is everything encrypted at rest?

General Questions

Q: What are Databricks AI-powered features?

Sharing: Freely share

Answer:

Databricks AI-powered features use a compound AI system that combines the use of AI models, retrieval, ranking, and personalization systems to understand your organization's data and usage patterns.

These AI-powered features enable everyone in an organization to be more productive using data and AI, while maintaining the governance and controls established in Unity Catalog.

Q: What is Genie Code?

Sharing: Freely share

Answer: Genie Code is your autonomous AI partner for data work. Genie Code accelerates complex data work by handling multistep tasks autonomously while adapting to your data and governance model.

Please see [\[Internal\] Genie Code FAQ](#) for more information.

Q: How do I enable Genie Code?

Sharing: Freely share

Answer: Genie Code is now GA and enabled by default, apart from in CSP workspaces (ring 1 customers – expected in April). You can follow the instructions ([AWS](#), [Azure](#), [GCP](#)) to disable certain AI features. Account admins have the ability to enable/disable the use of third party model partners and prevent workspace overrides using the **Enforce** toggle.

Genie Code is GA and we do not expect customers to disable it unless there is a compliance/legal issue. In that case, customers can follow the SOP [here](#).

Q: Is Genie Code enabled by default for workspaces?

Sharing: Freely share

Answer: Yes, most workspaces have Genie Code and any services that depend on 3P AI features are enabled by default. You can follow the instructions ([AWS](#), [Azure](#), [GCP](#)) to enable or disable certain features. Account admins have the ability to enable/disable the use of third party model partners and prevent workspace overrides using the **Enforce** toggle.

Q: What happens when I enable the [Compliance Security Profile](#) on my workspace? Will Genie Code be disabled automatically?

Sharing: Freely share

Answer: Genie Code and autocomplete are HIPAA and PCI compliant and auto-enabled in CSP HIPAA and PCI workspaces. All other features and CSP profiles will have partner-powered AI features disabled by default.

Q: Are there plans for more fine-grained access controls (per-workspace, per-user), or is the feature only available per account?

Sharing: Freely share

Answer: We do not currently have per-user feature access controls planned.

Q: If someone enables DatabricksAI-powered features for my account or workspace is that captured in my audit logs?

Sharing: Freely share

Answer: Yes. This action will be captured in the account-level and workspace audit events ([AWS](#), [Azure](#), [GCP](#)). For example, the following system table query will show you account and workspace enable/disable events for (at the time of writing) all AI-powered features:

```
None
with account_enablement AS (SELECT
  event_date,
  event_time,
  audit_level,
  workspace_id,
  source_ip_address,
  user_identity.email,
  action_name,
  request_params.settingTypeName,
```

```

request_params.settingValueForAudit,
response.status_code
FROM system.access.audit
WHERE service_name = 'accounts'
AND action_name = 'setSetting'
AND request_params.settingTypeName IN ('assistant', 'assistant_acct_enable',
'assistant_acct_allow_override')),
workspace_enablement AS (SELECT
event_date,
event_time,
audit_level,
workspace_id,
source_ip_address,
user_identity.email,
action_name,
request_params.settingTypeName,
request_params.settingValueForAudit,
response.status_code
FROM system.access.audit
WHERE service_name = 'workspace'
AND action_name = 'setSetting'
AND request_params.settingTypeName IN ('assistant_ws', 'enable_datarooms',
'ghost_text_workspace'))

SELECT * FROM account_enablement
UNION ALL SELECT * FROM workspace_enablement
ORDER BY event_time DESC

```

Q: Are the requests sent by a user to Genie Code captured in audit logs?

Sharing: Freely share

Answer: Databricks Assistant has its own system table for tracking user interactions with the Assistant ([AWS](#), [Azure](#), [GCP](#)). However, it does not include the actual requests sent by the user.

INTERNAL ONLY:

See [\[Mini-PRD\] System Table Support for Agent Mode Execution](#) for what's coming

Q: Are the questions sent by a user within an AI/BI Genie space captured in audit logs?

Sharing: Freely share

Answer: AI/BI Genie events are captured within Databricks audit logs ([AWS](#), [Azure](#), [GCP](#)). These do not include the actual questions sent by a user, although a user with CAN MANAGE permissions can use the History tab to view questions and responses unless they have been deleted by a user.

Q: What is Genie Code Autocomplete?

Sharing: Freely share

Answer: Genie Code Autocomplete is a sub-feature of Genie Code that provides inline code suggestions as you type for Python and SQL in notebooks, the SQL editor, and the File editor. It works with a Databricks-managed model to provide low-latency (~500ms) suggestions as you type.

Please see the announcement [blog post](#) for more information.

Q: How can I tell whether my customer has Partner-Powered AI enabled or disabled?

Sharing: Internal FAQ

We typically recommend asking [Databricks One](#) in Central Logfood – it's remarkably good at addressing this question.

Two dashboards:

- [LLMProxy Model Availability Settings](#) — per-workspace effective PPAI state, account/workspace values, enforce flag, and model availability
- [Genie Ready Review](#) — broader readiness (AIM / SCIM / UC + PPAI)

Model Questions

Q: What models are AI-powered features using?

Sharing: Only shared if requested

The best source for this information is our public doc pages ([AWS](#), [Azure](#), [GCP](#)).

Importantly, Databricks is following a multi-AI strategy - our approach is to move to the most performant and accurate (and also readily available) model(s). As such, the underlying models could change but the following foundational commitments apply:

1. Our trust and safety commitments ([AWS](#), [Azure](#), [GCP](#))
2. We provide transparency of our model partners ([AWS](#), [Azure](#), [GCP](#))
3. Any model partners are subject to our [Subprocessor](#) / [Subprocessor Changes](#) / [Subprocessor Objections](#) commitments
 - NB - for Azure Databricks the Microsoft Online Services Subprocessors commitments apply
4. Customers can self service opt-out of partner-powered AI via an account/workspace level control

Please ensure that the customer is aware of the above before sharing specific model details. Customers should expect models to evolve behind the scenes and without customer notification/warning.

- NB - depending on the scope of the change we may provide release note updates, and of course we cannot introduce a new subprocessor without honouring our contractual commitments - see above).

If your customer is requesting absolute transparency on which models are being used behind the scenes please raise go/brickroad issues to track this demand.

Feature	Models Used	Launch status	Governed by "Partner-powered AI" setting?
Genie Code Autocomplete	• Databricks-hosted fine-tuned OSS model ◦ <i>Currently the model is based on a FIM model from the Qwen family but telling the customer that may not necessarily help them.</i>	GA	• No
Genie Code	• Claude Opus 4.6	GA	• Yes, but falls back to

	○ Currently evaluating / AB testing Claude Sonnet 5		Databricks-hosted gpt-oss-120b
Genie Code Agent mode	● Claude Opus 4.6 ○ Currently evaluating / AB testing Claude Sonnet 5	GA	● Yes – feature is unavailable when disabled
Genie Agents	● GPT-4.1 ○ Currently evaluating / AB testing GPT-5.4	GA	● Yes – feature is unavailable when disabled
Genie Agents Agent mode	● Claude Sonnet 4.5 ○ Currently evaluating / AB testing Claude Sonnet 5	GA	● Yes – feature is unavailable when disabled
Genie One	● Claude Opus 4.6 ○ We actually use different models for different tasks (for example GPT-5.4-mini is used for table search/re-rank and GPT-5.6-Luna is used for web search). Currently Claude Opus 4.6 is the primary model and powers Genie One Chat.	GA	● Yes – feature is unavailable when disabled
AI-generated comments	● GPT-4o-mini	GA	● Yes but falls back to Databricks-hosted OSS model

Table last updated: Aug 6, 2026

INTERNAL ONLY:

- The table above is an over-simplified version. In reality we often have fallback models for different features, and may do A/B testing on a subset of traffic to trial new models (this A/B testing is performed in Databricks-owned environments like Dogfood first). [We're also seeing huge improvements using multiple models for the same task](#). Please see the comment above and help customers understand what they're using: a managed service with model provider transparency and a foundational set of commitments to ensure that customer data is protected.

Q: Are there plans to integrate with other models?

Sharing: Freely share

Answer: Customers can optionally configure the use of Genie Code with Databricks-hosted models by disabling the partner-powered AI features toggle, however not all features are supported and quality will be lower.

Q: Why doesn't Databricks support customer-provided models, like using a customer-provided OpenAI endpoint?

Sharing: Only shared if requested

Answer: AI features are managed "talk to your data" solutions, rather than frameworks that allow a customer to build their own models or their own solutions on top of models. We are evolving these features regularly to use the model that gives us the best quality and need to reserve the right to change that model as needed. Even when that model used is constant for a period of time, customers bringing their own models will inevitably lead to differences in quality and latency.

Q: What is the difference between OpenAI and the Azure OpenAI Service?

Sharing: Freely share

Answer: Genie Code uses the [Azure OpenAI](#) Service, which means that all data sent to the service is managed under the Azure terms and subscription. The [Azure OpenAI](#) Service is operated by Microsoft rather than by OpenAI.

Q: If you change the underlying models or services between providers will I be notified?

Sharing: Only shared if requested

Answer: In the event of changing the models or services between third party providers, we will notify customers via what's coming and release notes.

Q: Do you use the same Azure OpenAI subscription/service for all customers? How are my requests isolated from other customers?

Sharing: Only shared if requested

Answer: Databricks uses the same Azure subscription but different geographically-based Azure OpenAI services based on the geographic location of the workspace. Importantly, Databricks has opted into the exemption from abuse monitoring and human review program, under which Microsoft does not store any prompts and completions sent to the Azure OpenAI service. Furthermore, all requests to / responses from Genie Code are handled with the same strong isolation layers that apply to all control plane services.

Q: Do you use the same Anthropic service for all customers? How are my requests isolated from other customers?

Sharing: Only shared if requested

Answer: The Anthropic endpoints (hosted via Amazon Bedrock within the Databricks security perimeter) operate on shared infrastructure at the geographic level. Multiple customers may route through the same endpoints, but requests are never co-processed. Each prompt is independently authenticated, authorized, and handled in a stateless, logically isolated manner, with no cross-customer data sharing at any layer. Anthropic does not store prompts or completions submitted through these endpoints. We support encryption in transit and at rest: TLS 1.2+ for in-transit data and AES-256 for at-rest data.

Q: Do you use the same Databricks-managed model serving endpoints for all customers? How are my requests isolated from other customers?

Sharing: Only shared if requested

Answer: Databricks uses Geo-based model serving endpoints, and the model serving endpoint may vary by the model being served / feature being used but is ultimately shared

between customers. These model serving endpoints are stateless, protected through multiple layers of isolation and implement the following security controls to protect your data:

- Every customer request to Model Serving is logically isolated, authenticated, and authorized.
- Databricks Model Serving encrypts all data at rest (AES-256) and in transit (TLS 1.2+).

The same set of [trust and safety commitments](#) apply:

- Databricks does not use data, prompts, or responses submitted to or output from these features to train generative foundation models that Databricks makes available for use by third parties.
- Data is only sent when you interact with Databricks AI-powered features. Databricks sends your prompt, relevant values, table metadata, errors, as well as input code or queries to help return more relevant results.

Q: How accurate is Genie Code?

Sharing: Freely share

Answer: All AI models can make mistakes, misunderstand prompts, hallucinate answers, and introduce bias. You are fully responsible for the code you execute.

Q: Will Genie Code execute dangerous code?

Sharing: Freely share

Answer: In general, Genie Code will not automatically execute code on your behalf. Where this is not necessarily the case (such as via [Agent mode](#)), you have several options to remain in the driver's seat:

- Allow once: approve a single execution
- Always allow for this thread: streamline work within the current conversation. This resets when you press the "+" in the top-right corner of the panel.
- Always allow: give approval until you change the setting

Please note that all AI models are error-prone and can make mistakes, misunderstand prompts, hallucinate answers, and introduce bias. You are fully responsible for the code you execute.

Agent Mode does have some built-in guardrails to prevent dangerous code execution, such as the destruction of data (e.g., tables/files/schemas), changes to system permissions, calls to external websites, or other malicious impacts. That said, we still recommend carefully reviewing generated code, especially when it touches production data, important tables, or other sensitive operations.

Data Questions

Q: What data is being sent to the models?

Sharing: Freely share

Answer: The data sent to the models varies by feature, but importantly:

- Databricks uses only the data necessary to provide the service. Data is sent only when you interact with Databricks AI features. Databricks sends your prompt, relevant table metadata, errors, and input code or queries to help return more relevant results. Relevant values may also be included (such as prior results or data sampling).
- We do not send information that the current user **does not have access** to.
- You can also refer customers to the AI-powered features trust & safety page for the above ([AWS](#), [Azure](#), [GCP](#))

Q: Does the metadata sent to the models respect the user's Unity Catalog permissions?

Sharing: Freely share

Answer: Yes, all of the data sent to the model respects the user's Unity Catalog permissions (i.e. it does not send metadata relating to tables that the user does not have permission to see).

Q: Where is the data being stored?

Sharing: Freely share

Answer: Metadata relating to AI-powered features is stored in the control plane database. The control plane database is AES-256-bit encrypted, and customers who require control over the encryption key for scoped data can utilize our [customer-managed key](#) feature.

For features like Intelligent Search, Databricks maintains a search index over the customer's content, such as:

- Notebook metadata/content,
- SQL query metadata/content
- UC metadata
- Dashboard metadata
- Job metadata
- Model metadata

This search index is AES-256-bit encrypted. Note that for customers who have enabled a [customer-managed key](#), we do not index data that we commit to encrypting with their key (such as the notebook / SQL query content mentioned above). In the future, we plan to extend CMK coverage to vector indices to include this.

For AI/BI Genie, user questions and Genie natural language responses are stored within the control plane database, which is AES-256-bit encrypted, and customers who require control over the encryption key can utilize our [customer-managed key](#) feature. Query results (i.e., the customer's row-level data) are not stored or cached in the control plane. These are cached in the SQL result cache (see the docs for [AWS](#), [Azure](#), [GCP](#)) and retrieved or recomputed when needed.

With Genie data sampling ([AWS](#), [Azure](#), [GCP](#)), customers can improve accuracy in a Genie space by allowing Genie to use sample data, helping it better match user-submitted prompts. Data sampling is enabled (excluding CSP workspaces), although users can select which columns should be sampled or disable it altogether for a space. Genie samples up to 255 distinct values per column, and each sampled value is truncated at 127 characters. When a column or row exceeds that maximum, only a subset is used. The computed data samples are stored in the customer's workspace storage account, subject to the encryption configured there (including the use of a [customer-managed key](#)). Sampled data can be provided alongside available metadata, comments, instructions, and queries to improve interpretation and query accuracy.

Note:

- Chat history is stored in the same place as other notebook content and is therefore encrypted via the CMK (if configured)
- Genie spaces created after April 10, 2025, are now encrypted with CMK if configured at the time of creation
- Like other workspace objects, the retention period for a chat thread or Genie space is scoped to the lifecycle of the object itself. If a user deletes the thread, then it will be deleted within 30 days

- Each user has their own chat history, which can get large if you have a long-running conversation, culminating in `Notebook size exceeds limit` errors. The solution right now is to clear long-running threads
- One other side effect of each user having their own chat history is that only you can see your prompts/responses

Q: Is my data being used to train models?

Sharing: Freely share

Answer: Databricks does not use data, prompts, or responses submitted to or output from these features to train generative foundation models that Databricks makes available for use by third parties.

Please refer to our [Databricks AI features trust and safety page](#) for more information.

INTERNAL ONLY:

- Databricks currently stores the inputs and outputs (i.e., requests and responses) as part of our service logs, but we do not use that data for model training and do not conduct any human review.
- Please note that we don't usually discuss our service logs with customers, and our approach to Genie Code isn't different from other parts of the platform.
- Our service logs are stored in the regional logfood associated with that workspace. They are AES-256-bit encrypted at rest, have restricted access policies, and a limited retention period.

Q: If I execute a query with results, and then ask a question, do the results of my query get sent to the model?

Sharing: Freely share

Answer: This is not generally the case, however:

- With Genie data sampling, customers can improve accuracy in a Genie space by allowing Genie to use sample data, helping it better match user-submitted prompts. Data sampling is enabled by default, although users can select which columns should be sampled or disable it altogether. Please refer to the docs ([AWS](#), [Azure](#), [GCP](#)) for more information.
- Agent Mode can also analyze cell outputs and read data samples from tables.

Q: Can I opt out of Azure OpenAI and Anthropic on Bedrock storing my data?

Sharing: Only shared if requested

Answer: Yes. Databricks has already opted into the [exemption from abuse monitoring and human review](#) programs, under which Microsoft does not store any prompts and completions sent to the Azure OpenAI service.

When Databricks features use Anthropic models, requests go to Databricks-managed endpoints for Anthropic hosted in our private tenant on Amazon Bedrock, within the Databricks security perimeter. Anthropic does not retain prompts or completions submitted through these partner-powered features and does not use this data to train their models.

Q: Can I request an exemption from abuse monitoring and human review on Azure OpenAI and/or Anthropic on Bedrock?

Sharing: Only shared if requested

Answer: Databricks has already opted into this exemption for Azure OpenAI Service. Anthropic follows the same geo-aligned, stateless, and logically isolated serving model, with additional safety hardening. See the [Claude Opus 4.6 Model Card](#) and the [Claude Sonnet 4.5 Model Card](#) for details.

INTERNAL ONLY:

- Databricks does not have an opt-out for abuse monitoring and human review for any logs stored on our side.

Q: Is any content filtering applied to the responses to filter for harmful responses?

Sharing: Only shared if requested

Answer: Yes. Genie Code leverages Azure OpenAI's synchronous content filtering set at the [default Medium, High filtering](#) level. Jailbreak is also on by default now. For Anthropic on

Bedrock, the Genie Code relies on Anthropic's built-in safety mechanisms and additional hardening against harmful outputs as described in Anthropic's safety documentation. See the [Claude Opus 4.6 Model Card](#) and the [Claude Sonnet 4.5 Model Card](#) for details.

INTERNAL ONLY:

- We can't enable asynchronous content filtering since data submitted by Azure Databricks in these features is not retained by Azure OpenAI, even for abuse monitoring. These features have been opted-out of Azure OpenAI's retention of data for that purpose.

Q: Are there any guardrails in place to strip out PII/sensitive data before sending it to Databricks-managed models or our model partners?

Sharing: Only shared if requested

Answer: No, we do not have mechanisms to strip out PII or sensitive data. It is up to the user to validate what data they are comfortable sending.

Q: What telemetry is captured?

Sharing: Only shared if requested

Answer: Databricks has opted into the [exemption from abuse monitoring and human review](#) program, under which Microsoft does not store any prompts and completions sent to the Azure OpenAI service. For Anthropic on Bedrock, requests are routed through Databricks-managed endpoints with zero data retention and no training by Anthropic on customer prompts or completions. Like most SaaS services, Databricks captures telemetry, which is used to monitor the platform's usage and health and support customers. We encrypt everything in our account.

INTERNAL ONLY:

- Databricks currently stores the inputs and outputs (i.e., the request to and response from Azure OpenAI) as part of our service logs, but Databricks does not use data, prompts, or responses submitted to or output from these features to train generative foundation models that Databricks makes available for use by third parties.

- Please note that we don't usually discuss our service logs with customers, and our approach to Genie Code isn't different from other parts of the platform.
- Our service logs are stored in the regional logfood associated with that workspace, they are AES-256 bit encrypted at rest, have restricted access policies and a limited retention period.
- The statement for Databricks-managed models comes from our model serving docs ([AWS](#), [Azure](#)). Note that at the time of writing it is forward looking – we're not currently doing this but may start to at some point in the future.

Q: Who owns the output data? If AI-powered features are generating some code, who owns that IP?

Sharing: Freely share

Answer: The customer owns their own output.

Q: Will AI-powered features generate copyrighted code? Am I legally protected if I use that in my production software?

Sharing: Only shared if requested

Answer: Databricks takes a number of measures to prevent AI-powered features from outputting harmful content, including code that may be protected by copyright. Moreover, Databricks promises to protect customers who have agreed to our [Service Specific Terms](#) ("SSTs") against certain claims that their use of AI-powered features infringes a third party's intellectual property rights, as outlined in the Databricks GenAI Indemnity section of our SSTs.

Legal & Compliance Questions

Q: Are there any data residency considerations I should be aware of?

Sharing: Only shared if requested

Answer: Databricks AI-powered features are Designated Services and comply with data residency boundaries at a Geo level.

Please see our AI-powered features trust and safety ([AWS](#), [Azure](#), [GCP](#)), Geos ([AWS](#), [Azure](#), [GCP](#)) and Designated Services pages ([AWS](#), [Azure](#), [GCP](#)) for more information.

We will continue to evaluate support for other regions for future versions to meet latency and data residency requirements.

INTERNAL ONLY:

- If your customer requires a geo/region that is not currently supported, please raise a [go/productblocker](#).
- Preview versions of the upcoming doc changes related to the above are available below:
 - AWS: [Databricks Geos: AI data residency](#)
 - AWS: [Databricks Designated Services](#)
 - Azure: [Databricks Designated Services](#)
- ***I'm not a lawyer comment:*** Geo-based data processing is aligned with GDPR. Under GDPR you can freely transfer data within the EU and to countries with an [adequacy agreement](#) (i.e. the UK, Switzerland)
- ***I'm not a lawyer comment:*** In actual fact, since [Databricks complies with the EU-U.S. Data Privacy Framework](#) based on GDPR [you can freely transfer data to Databricks in the US](#).
- Note that due to things like GPU availability, Geo (rather than region) based data processing commitments for GenAI features is very common:
 - [Azure copilot](#)
 - [Azure copilot for Fabric](#)
 - [Azure OpenAI](#)
 - [Amazon Q Developer](#) (Amazon Q Business is only [available in the US](#))
 - [Amazon Bedrock](#)
 - [Google Gemini](#)

Q: Can I request for my data to be deleted?

Sharing: Only shared if requested

Answer: Databricks automatically deletes usage logs after 30 days per our log retention policy. We do not support manual deletion of customer logs prior to 30 days.

Configurations (such as notebooks) stored within the Databricks application can be deleted

by customers via the UI or API as needed. After termination of a Databricks workspace, the customer content contained therein is deleted within 30 days of termination.

Q: Can I use Genie Code with tables that contain regulated data (FEDRAMP/PHI/PCI/IRAP)?

Sharing: Freely share

Answer: Yes, but if you want to process regulated data you must have the Compliance Security Profile ([AWS](#), [Azure](#)) enabled.

Azure Compliance Offerings PDF can be downloaded from [here](#). From this document, we can tell that Azure OpenAI is PCI and HIPAA compliant. The security team also did a significant audit of Azure OpenAI last year and determined it was Fedramp authorized. Therefore, Genie Code is covered by all of our existing Commercial certifications and Fedramp Moderate.

Internal Only

Mar 18, 2024 : Assistant received sign-off from the compliance and security team:
<https://databricks.atlassian.net/browse/SASS-3563>

Q: Can I enable/use Genie Code in Compliance security profile enabled workspaces?

Sharing: Freely share

Answer: Genie Code can be enabled and used in Compliance security profile ([AWS](#), [Azure](#)) enabled workspaces. It is up to customers to decide whether to use Genie Code, weighing up the assurances provided against the sensitivity of the data that they are likely to send to it.

Q: For AWS GovCloud and Azure Government customers, do we use Azure OpenAI and/or Anthropic instances in Azure Government?

Sharing: Only shared if requested

Answer: [As per the docs](#), Genie Code Agent mode requests on AWS GovCloud workspaces are routed to Azure OpenAI endpoints in Azure Government for LLM inference. Customer content does not leave U.S. government cloud regions.

For other Genie features, deployments follow available commercial cloud infrastructure; government cloud availability depends on model support and regional service maturity.

Q: What legal terms govern the use of AI-powered features? Is it the Azure OpenAI API terms or other terms determined by Databricks?

Sharing: Only shared if requested

Answer: Usage is governed by the existing Databricks terms and conditions the customer has agreed to when using Databricks.

Q: Does Databricks SOC2 report cover AI-powered features?

Sharing: Only shared if requested

Answer: Yes. The Databricks SOC2 covers the Databricks Data Intelligence Platform.

Q: Has there been a penetration test for AI-powered features?

Sharing: Only shared if requested

Answer: AI-powered features are implicitly included in our platform penetration test scope. Depending on the feature, many of the components that might be used (such as our serverless model serving architecture) have been subject to an independent penetration test. There are no plans for a feature-specific penetration test of AI-powered features.

Q: As part of the [Microsoft Responsible AI](#) review, what harm categories were assessed?

Sharing: Only shared if

requested

Answer: The following harm categories were assessed

- Harmful content
 - Hate and Fairness
 - Sexual
 - Self-harm
 - Violence
- Advice related to highly regulated domains, such as medical, financial, and legal
- 3rd party content regurgitation
- Groundedness (Fabrication and ungrounded content)
- Prompt injection (jailbreaks)
- Generation of malware, incorrect or insecure code

Encryption Questions

Q: Is everything encrypted in transit? Does any traffic hit the public internet?

Sharing: Freely share

Answer: For Databricks-managed models on AWS, traffic stays on the AWS network and does not use the public internet.

For Databricks-managed models on Azure, traffic stays on the cloud providers network and does not use the public internet. All traffic to Azure OpenAI services stays on the Azure network too.

INTERNAL ONLY:

To secure this traffic:

- Databricks performs TLS certificate verification.
- Databricks explicitly disallows modes and ciphers that are known to be insecure.

Q: Is everything encrypted at rest?

Sharing: Freely share

Answer: Databricks has opted into the [exemption from abuse monitoring and human review](#) program, under which Microsoft does not store any prompts and completions sent to the Azure OpenAI service. For Anthropic on Bedrock, Databricks enforces zero data retention and does not permit training on customer prompts or completions by Anthropic. Any data stored within a Databricks workspace is AES-256 bit encrypted. Customers that require control over the encryption key can utilize the Customer-Managed Key feature.

INTERNAL ONLY:

- Databricks currently stores the inputs and outputs (i.e. the request to and response from Azure OpenAI) as part of our service logs, but Databricks does not use data, prompts, or responses submitted to or output from these features to train generative foundation models that Databricks makes available for use by third parties.
- Please note that we don't usually discuss our service logs with customers, and our approach to Genie Code isn't different from other parts of the platform.
- Our service logs are stored in the regional logfood associated with that workspace, they are AES-256 bit encrypted at rest, have restricted access policies and a limited retention period.