

Databricks Foundation Model APIs, External LLMs, and Genie

Customer security and data-residency briefing

Prepared: August 2026

Introduction

Databricks is designed to let enterprises use generative AI without giving up control of governed data, model access, or geographic processing boundaries.

The key points are:

- Databricks Foundation Model APIs and Databricks AI assistive features, including Genie experiences, are Designated Services. Databricks Geos determine where customer content is processed. When the service and selected model are available in the workspace Geo, processing remains in that Geo. For a strict European boundary, customers should enable the workspace setting that enforces processing within the workspace Geography for Designated Services and select only models/features documented as available in that Geo.
- European data is not routed to the United States or Asia when Geo enforcement is enabled and the selected service/model is available within the European Geo. If cross-Geo processing is enabled, or a model is unavailable in the workspace Geo, content may be processed in another permitted Geo. This is why model-by-model and cloud-by-cloud verification is important.
- Databricks does not use customer prompts, responses, or submitted data to train the generative foundation models Databricks makes available to third parties. For paid Model Serving accounts, Databricks also states that inputs and outputs are not used to train models or improve Databricks services.
- Partner-powered Databricks AI features use zero-data-retention endpoints from model partners. In practical terms, partner providers do not retain the prompts or responses submitted through those features. This commitment is separate from Databricks service telemetry and feature storage, which must be reviewed by feature.
- Unity Catalog remains the governing layer for data and AI workflows. Genie and other AI features only send model context that the requesting user is authorized to access, and data access remains subject to Unity Catalog permissions, row filters, column masks, and related controls.

These statements describe the published product posture. A production approval should additionally validate the specific cloud, workspace region, feature, model, serving mode, compliance profile, cross-Geo setting, and applicable model terms.

1. Databricks Foundation Model APIs

Foundation Model APIs provide workspace-level access to Databricks-hosted foundation models through Model Serving. The service supports pay-per-token and provisioned-throughput modes, with model and region availability varying over time.

Security controls

- Requests are authenticated and authorized as workspace API requests, with workspace access controls and network controls such as IP allowlists and supported Private Link configurations.
- Requests are logically isolated and stateless. The serving endpoint does not persist the request context as model memory.
- Traffic is encrypted in transit using TLS 1.2 or higher, and Databricks-stored data is encrypted at rest using AES-256 or stronger controls where applicable.
- Administrators can restrict which Databricks-hosted foundation models users may invoke through Unity Catalog permissions for Foundation Model APIs.

Residency controls

Foundation Model APIs are a Databricks Designated Service and use Databricks Geos to manage processing geography. The authoritative controls are:

- the workspace Geo;
- the availability table for the selected model and serving mode;
- the account/workspace setting that enforces processing within the workspace Geography for Designated Services; and
- whether cross-Geo processing is enabled.

The Designated Services documentation states that, when a Designated Service is available in the workspace Geo, customer content is processed in that Geo. If cross-Geo processing is enabled, a service may process content in another Geo where required for availability. The same documentation also states that customer data is not stored outside the workspace Geo merely because cross-Geo processing is enabled.

For an EU-residency-sensitive workload, the recommended control pattern is:

- keep the workspace in the applicable European Geo or EU Data Boundary;
- enable enforcement of processing within the workspace Geography for Designated Services;
- choose a model listed as available in that Geo; and
- test the configuration before production use, including failover and model-version changes.

Retention and training nuance

Databricks states that Foundation Model API inputs and outputs may be temporarily processed and stored for abuse prevention, detection, and mitigation, isolated from other customers and stored in the workspace region for up to 30 days. This is operational security handling, not model training.

Partner-provider retention can also have narrow safety exceptions for particular model classes or use cases. The current Model Serving documentation identifies model-specific exceptions, including safety retention requirements for certain Anthropic model classes and possible OpenAI safety retention for specified coding/routing scenarios. These exceptions should be checked against the exact model selected.

2. External Models: OpenAI, Anthropic, Gemini, and other providers

External Models are third-party models hosted outside Databricks and configured as Model Serving endpoints. They are different from Databricks-hosted Foundation Model APIs.

Databricks provides a centralized control point for external endpoints, including:

- endpoint-level access control and workspace authentication;
- centralized credential management using Databricks secret scopes;
- rate limits and usage controls; and
- a common interface for calling different model providers.

The customer remains responsible for the external provider configuration and contractual terms. Before approving an External Model, the customer should verify the provider's processing region, retention policy, training policy, abuse-monitoring policy, subprocessors, encryption, and deletion commitments. Data sent to an External Model is sent to the endpoint and provider that the customer configured; Databricks governance does not convert a customer-managed external endpoint into a Databricks-hosted model.

Provider-specific decision points

- OpenAI: verify the exact API or hosted offering, region, applicable enterprise/API data terms, and any safety-retention exception. Do not assume that every OpenAI route has identical retention behavior.
- Anthropic: verify whether the customer is using a Databricks-hosted/private offering or a customer-configured Anthropic endpoint. Check model-specific safety-retention terms, particularly for model classes that impose safety retention.
- Gemini: verify whether the model is available as a Databricks-hosted offering or is being called as a customer-configured external endpoint. Review the current Databricks model terms and Google/Gemini data-handling terms for the selected model.

For the lowest operational complexity, customers commonly prefer Databricks-hosted model offerings where the published Databricks controls, Geo behavior, and model terms can be assessed as one integrated service. For a customer-managed External Model, the provider contract and configuration remain a material part of the security boundary.

3. Databricks Genie and Unity Catalog governance

Genie is a governed natural-language interface over selected enterprise data. The Genie family includes Genie One, Genie Agents, and Genie Code. The experience is not an unrestricted chatbot with a free-standing copy of the data.

What Genie sends to models

Databricks documents that AI assistive features send only the data needed to provide the feature. Depending on the experience, this can include:

- the user's natural-language prompt;

- table and column names and descriptions;
- relevant values or configured sample values;
- example SQL, functions, instructions, and prior questions; and
- errors, code, query context, or cell output for certain Genie Code modes.

Unity Catalog enforcement

Databricks documents that all data sent to AI assistive feature models respects Unity Catalog permissions. A user cannot use Genie to bypass access to data they cannot access through the governed data layer.

For enterprise governance, Unity Catalog provides the control surface for:

- catalog, schema, table, view, function, volume, and model privileges;
- row-level filtering and column masking where configured;
- attribute-based access control and governed tags where available;
- lineage and auditability for governed data and AI assets; and
- access to the SQL warehouse and Genie space itself.

Genie Agents are designed for read-only interaction with governed data and generate/run read-only SQL. Genie Code has broader agentic capabilities in supported modes, so tool approval, workspace permissions, and operational controls should be reviewed separately for developer workflows.

Storage and visibility

Genie conversations, responses, and related application data can be stored in Databricks-managed control-plane or workspace storage depending on the feature. Databricks documents encryption at rest and customer-managed-key options for applicable managed services. Partner model providers do not retain prompts or responses when partner models are used through Databricks AI assistive features, but this does not mean that a Genie conversation is never stored inside Databricks. Customers should review chat sharing, workspace retention, audit logging, and feature-specific storage behavior as part of their data classification process.

4. Zero data retention: what it means and what it does not mean

Zero data retention is a partner-endpoint control. It means that the model partner does not retain the prompt or response after processing, including for ordinary abuse-monitoring workflows, for the covered Databricks AI feature path.

It does not mean:

- no request is ever transmitted to a model;
- Databricks never stores feature data, chat history, audit data, metrics, or security telemetry;
- every provider, model family, or customer-configured External Model has identical retention terms; or
- model-specific safety-retention requirements are overridden.

The most accurate customer statement is therefore: Databricks uses zero-data-retention partner endpoints for the covered partner-powered AI features, does not train base/generative foundation

models on customer submissions, and documents separate Databricks storage and model-specific safety-retention conditions that must be evaluated for the selected feature and model.

5. Customer verification checklist

Before production approval, the customer's security and privacy teams should record:

- cloud provider and workspace region;
- Databricks Geo and applicable EU/UK or other data-boundary requirement;
- whether "Enforce data processing within workspace Geography for Designated Services" is enabled;
- whether cross-Geo processing is enabled;
- exact Genie experience and exact model/provider used;
- whether the model is Databricks-hosted, partner-powered inside Databricks, or customer-configured External Model;
- data categories, including personal data, confidential data, regulated data, and sample values that may be sent to the model;
- Unity Catalog privileges, row filters, column masks, Genie space ACLs, and SQL warehouse permissions;
- chat, prompt, response, audit, and service-log retention settings;
- applicable Databricks service terms, DPA, subprocessors, and model-specific terms; and
- a change-management trigger for model substitutions, new model versions, new regions, or new Genie capabilities.

6. Public-facing resources for customer verification

Databricks Security and Trust Center

- [Databricks Security and Trust Center](#) – central portal for security, privacy, compliance, architecture, and trust materials.
- [Databricks AI Security](#) – AI security risks and recommended controls across the AI lifecycle.
- [Databricks Responsible AI](#) – responsible AI principles, governance, and safety practices.
- [Databricks AI Security Framework \(DASF\)](#) – public framework for assessing AI risks and controls.
- [Announcing the Databricks AI Security Framework 2.0](#) – blog overview of the DASF 2.0 and its operationalization resources.
- [Databricks Subprocessors](#) – current public subprocessor and provider reference.
- [Databricks Privacy Notice](#) – general privacy notice.

Residency and Designated Services

- [Databricks Geos: Data residency – AWS](#) – Geo definitions, processing boundaries, and cross-Geo processing.
- [Databricks Designated Services – AWS](#) – Geo-by-Geo availability matrix for Foundation Model APIs and Genie-related AI features.
- [Databricks Geos: Data residency – Azure](#) – Azure Geography behavior and enforcement controls.
- [Databricks Designated Services – Azure](#) – Azure Geography availability and the cross-Geo caveat.

Foundation Model APIs and Model Serving

- [Databricks Foundation Model APIs – AWS](#) – overview, serving modes, and supported usage patterns.
- [Foundation Model APIs compliance and security – AWS](#) – processing geography, regional model availability, access control, and network security.
- [Foundation Model APIs compliance and security – Azure](#) – Azure-specific residency and compliance details.
- [Deploy models using Model Serving – AWS](#) – data protection, encryption, training, retention, and partner-provider caveats.
- [Databricks-hosted foundation models – Azure](#) – model list and model-specific notes.
- [Applicable model terms – AWS](#) – current model terms and provider policy links.
- [External models in Model Serving](#) – centralized governance and credential management for third-party models.
- [ADI Services](#) – public documentation for Databricks-hosted OpenAI and Gemini offerings and related model terms.

Genie and Unity Catalog

- [Databricks AI assistive features trust and safety – AWS](#) – data sent to models, no-training commitment, partner retention, encryption, permissions, and residency.
- [Databricks AI assistive features trust and safety – Azure](#) – Azure-specific trust and safety details.
- [Genie](#) – product family overview.
- [Genie Agents](#) – governed natural-language access to selected datasets.
- [Create and manage a Genie Agent](#) – configuration, datasets, instructions, and trusted-answer setup.
- [Genie Code](#) – developer and data assistant capabilities and Unity Catalog governance.
- [Unity Catalog access control](#) – privileges and access-control model for governed data assets.
- [Unity Catalog](#) – public product overview for unified data and AI governance.

Blogs, videos, and training

- [Governing coding agent sprawl with Unity AI Gateway](#) – public blog on centralized security, audit, and model governance.
- [Introducing new governance capabilities to scale AI agents with confidence](#) – public blog on model, tool, and data governance.
- [Data Intelligence with AI/BI Genie and Azure AI Foundry](#) – on-demand Genie and governance demonstration.
- [What is AI/BI Genie?](#) – public Genie overview and data-to-SQL flow.
- [AI-Powered Data Governance & Insights with Azure Databricks](#) – demonstration of Unity Catalog permissions, lineage, and Genie.
- [Databricks Training and Certification catalog](#) – Gen AI Fundamentals, Data Management and Governance with Unity Catalog, and Generative AI Application Evaluation and Governance learning paths.
- [Official Databricks YouTube channel](#) – current product, AI, governance, and security videos.

Conclusion

Databricks' security model is based on governed data access, model optionality, encrypted and isolated serving, Geo-aware processing, and explicit controls over partner-powered AI. For a high-sensitivity workload, the right approval question is not simply "Does Databricks use an LLM?" It is "Which feature, which model, which Geo, which administrative settings, which data is sent, and which retention terms apply to this exact workflow?"

Databricks' public documentation provided [here](#) to make those answers verifiable.

Use the links above as the baseline evidence set and retain a configuration-specific assessment for the production workspace.