



Databricks

Unity AI Gateway



August 2026


Agenda for this session

- 1 Unity AI Gateway
- 2 Unity AI Gateway Deep dive
- 3 Omnigent
- 4 Q&A

Agenda for this session

- 1 Unity AI Gateway
- 2 Unity AI Gateway Deep dive
- 3 Omnigent
- 4 Q&A


Today, everyone is building with AI





There is
agent sprawl

Sprawl we all see: Coding Agents



Agent sprawl turns simple questions into governance issue across the C-suite



Head of AI

How many agents are running, and who owns them?



CISO

Who can use an agent?



Chief Data Officer

What data did the agent access?



Chief Risk Officer

How can I keep an agent from deleting a data set?



AI Platform Admin

How should we monitor quality and debug issues?



CFO

What does each agent cost, and how much did we spend?






Unity AI Gateway

Centralized governance, security, and cost controls over agents, models, and tools

Unity AI Gateway

Control and audit your entire AI estate
with unified governance, visibility, spend
controls, and guardrails



✓ **Open governance for multi-AI**
Govern frontier and OSS models, agents, MCPs, and third-party AI frameworks consistently across Databricks and external platforms

✓ **Govern access and AI spend**
Control models, agents, and MCPs with centralized policies, budgets, rate limits, and identity-aware controls


✓ **End-to-end AI observability**
Capture prompts, traces, tool calls, payload logs, and policy decisions across every AI interaction





Unified AI tokenomics and spend governance

Your single source of truth for AI
consumption and spend



Multi-AI spend visibility

Track AI costs across Databricks and external providers



Contextual budget policies

Allocate budgets by team, application, agent, or coding tool



Hard spend caps

Automatically stop requests before budgets are exceeded





Unity AI Gateway

Agent Registry

Access Control

Contextual Policies

Budgets

Smart Routing

Agent Tracing

Models




Agents




MCP & Skills



External Agents




AI Coding



**Unified
governance**
for your entire
AI estate

Unity AI Gateway has agents, tools, and skills



Agent Registry

Govern with fine-grained access controls

Discover all
agents, including
external agents

Unity AI Gateway

Govern and access your AI. [Learn more](#)


Models **Beta** Providers **Beta** MCPs Skills **Agents**

Search

Filter by Agent Type

Name	Agent type	Owner
retail_agent	Custom Agent	elise@acme
sec_10k_parser	Document Intelligence	archika@acme
opsbot	Custom Agent	adam@acme
aws_agentcore_supp...	External Agent	david@acme
finance_langgraph_a...	Custom Agent	jacqueline@acme
smart_router	Custom Agent	bilal@acme
RL_agent	Custom Agent	tejas@acme


Contextual security policies




Agent Registry

Contextual
Policies

Define in SQL custom security policies
Dynamic, stateful, and contextual to data




Contextual security policies




Agent Registry

Contextual
Policies

Define in SQL custom security policies
Dynamic, stateful, and contextual to data



Integrate policies with an open ecosystem



Agent Registry

Contextual
Policies


Connect to external partner guardrails for governance
Integrate pre-built policies into your agents

ALICE



+ many more partners

Monitor costs and set budgets



Agent Registry

Contextual Policies

Budgets

Total Policies

5

Exceeded

2

5 users blocked

With Enforcement

2

blocking or throttling

🔍 Search budgets...

NAME	WORKSPACES	TAGS	USERS & GROUPS	STATUS	WHEN EXCEEDED
default-monthly-budget	All	—	all-users per-user · monthly	4 of 11 exhausted	<button>🚫 Block Access ▾</button>
power-user-budget	All	—	power-users per-user · monthly	1 of 4 exhausted	<button>🚫 Block Access ▾</button>
prod-design-budget	All	—	product-design shared · weekly	15% spent	<button>⚠️ Warn Only ▾</button>
dev-gateway-budget	development	—	—	38%	<button>⚠️ Warn Only ▾</button>
chatbot-weekly-budget	All	chatbot	—	52%	<button>⚠️ Warn Only ▾</button>

Smart Routing

 Unity AI Gateway


Agent Registry

Contextual Policies


Budgets

Smart Routing


Intelligent routing



Cross-provider and cloud fallbacks
& traffic splitting



Agent Tracing



Agent Registry

Contextual Policies

Budgets

Smart Routing


Agent Tracing

Analyze coding logs to
spot bad practices

Q1: Are People Using the Wrong Model for the Task?

Yes — almost everything defaults to Opus regardless of complexity. 82% of sessions (134/163) use claude-opus-4-8. There is no model-routing behavior.

Q1: Model Mismatch — Where Does the Spend Go?



Integrate with Lakewatch
for alerts & monitoring

Unity AI Gateway > Governance

Governance


Budgets Traces Incidents

Severity	Incident	State	Created
Medium	sec-7861 PII detected in LLM prompt i...	Open	Jun 14, 2026,
Medium	sec-7860 PII detected in LLM prompt i...	Open	Jun 13, 2026,
Medium	sec-7859 PII detected in LLM prompt i...	Open	Jun 13, 2026,
Medium	sec-7858 PII detected in LLM prompt i...	Open	Jun 13, 2026,

End-to-end visibility into every call, cost, and decision

From Tokenmaxxing to Valuemaxxing

Monthly AI spend



Governing coding tools with Unity AI Gateway

Leverage every coding tool with any model with **centralized** governance

[Unity AI Gateway](#) > Governance

Governance

Budgets **Traces** Incidents

Time range: Last 24 hours ▾

unity_gateway_otel_spans ▾

vlatko-test ▾

Search traces by id, request, or...

Filters ▾

Sort: Time ▾


Columns 4/6 ▾

100 of many

Time	Trace ID	Service	Principal
06/16/2026, 21:55:20	9c364490cba24a059a5e3ae86e86a...	LLM databricks-gemini-3-5-flash	ankit.mathur@databricks.com
06/16/2026, 21:55:19	9b84c55444e7443986247eef4307d...	LLM databricks-claude-opus-4-8	ankit.mathur@databricks.com
06/16/2026, 21:55:18	59cac9e2ca0143519933a91fab292161	LLM databricks-gemini-3-5-flash	ankit.mathur@databricks.com
06/16/2026, 21:55:16	68b213a8689048c7bb8e73eb3c49...	LLM databricks-gemini-3-5-flash	ankit.mathur@databricks.com
06/16/2026, 21:55:15	b393118c3dcd4ea1901e0f773110cd82	LLM databricks-gpt-5-4-mini	ankit.mathur@databricks.com
06/16/2026, 21:55:14	1b7e60642d424583b12c3bb6bacc2...	LLM databricks-gpt-5-4-mini	ankit.mathur@databricks.com

Analyze coding traces

Model Selection Verdict Breakdown: Actual vs Optimal Cost



Key Findings

- **OVERKILL (401 sessions, 98.3%)** — The vast majority of sessions used a model far more powerful than needed. The most commonly recommended alternative is **claude-haiku-4-5**, which would save **\$844.01 (90.2%)** of the cost in this category. [2]
- **APPROPRIATE (6 sessions, 1.5%)** — Only 6 sessions were correctly matched to their model. These used **claude-opus-4** appropriately, with zero potential savings. [2]

Ask a question, @ for objects, / for skills...




Always review the accuracy of responses.

Analyze coding traces

What Were People Doing?

OVERKILL Sessions by Task Category



The task categories break down as:

1. **Simple/straightforward tasks** (148 sessions, \$329 in savings) — Standard coding tasks like creating PRs, resolving PR comments, setting up dev environments, and making routine code changes. These don't require frontier-level reasoning.



Unity AI Gateway

Access
control

Agent Registry

Contextual Policies

Budgets

Smart Routing

Agent Monitoring

Models




Agents




MCP & Skills



External Agents



AI Coding



**Unified
governance
for your entire
AI estate**




AI assets are Unity Catalog securables

Same governance you built for data now covers models, MCPs, and agents.


Register once, govern it centrally: Each AI asset is a UC "service" securable (catalog → schema → service): Models, MCP, Skills, Agents

RBAC: grant EXECUTE (use), MANAGE (administer), CREATE SERVICE / CREATE CONNECTION (register)

Unified Governance: One governance model that extends to every next asset type. Same discovery, lineage, and audit as data.




Unity AI Gateway Architecture: How it Works?



Unity AI Gateway: How it works?

Every request takes the same governed path



Where is everything recorded?



Usage

system.ai_gateway.usage

Analyze request patterns (volume, latency, token usage)

Understand which models are being called and by whom

Monitor guardrail activity

Track token consumption at the request level

Distinguish between API types (e.g., ai_query vs. other)

Answer: "Which models are most popular?" or "What's the average latency per endpoint?"

system.billing.usage

Understand costs and DBU consumption for AI Gateway (or any other product)

Track spending trends over time

Allocate costs by workspace, SKU, or custom tags

Answer: "How much are we spending on AI Gateway?"



Telemetry

MLflow traces (mlflow_experiment_trace_otel_spans, _logs, _metrics in your catalog/schema)

Structured execution graph (spans with inputs, outputs, timing, attributes)

in OpenTelemetry (OTEL) standard format

Inference Tables (user specified table in UC)

Raw request/response JSON payloads

Being deprecated in favour of MLflow traces




Integración Bidireccional

DIRECCIÓN	PATRÓN	STATUS
Foundry → DBX Data	RunNow / RunCommand / Statement Execution via SDK	Available
Foundry → DBX AI	Genie API (databricks-langchain), Model Serving	GA
DBX → Foundry (Inference)	Register Foundry models en AI Gateway	Available (Jun 2026)
DBX → Foundry (Eval)	Databricks Evals para evaluar modelos/agentes Foundry	Available
Foundry Agent → Genie	Native Connector en AI Foundry (Build May 2025)	GA



Native Connector: **Genie** in AI Foundry

Announced at Microsoft Build — May 2025



Auth: OAuth U2M + OBO

1 Usuario en Teams → Copilot Studio agent

2 Redirect → Databricks OIDC authorize (PKCE)

3 SSO via Entra ID + consent screen

4 Auth code → access_token + refresh_token

5 MCP/Gateway invocado con access_token

6 Unity Catalog valida user_identity + GRANT

7 SQL ejecutado as user → resultados devueltos

OBO

Backend MCP calls propagan identidad del usuario

Audit

system.access.audit captura user_identity = end user



Agent Bricks vs Foundry Agent Service

⚡ Competitivo — coexistencia solo vía Gateway-fronts-Foundry

DIMENSIÓN	AGENT BRICKS	FOUNDRY AGENT SVC
Ecosystem	Databricks-native, cross-cloud	Microsoft-native, Azure only
Governance	Full UC (data + AI)	Azure RBAC / Entra
Tracing	MLflow 3 (open)	Azure Monitor
Orchestration	Supervisor multi-agent	Agent Framework 1.0
MCP	MCP-first architecture	Tool use / functions
Models	Any via AI Gateway	Azure OpenAI primarily
Surfaces	Any (APIs, Review App)	Teams, Copilot Studio



Agenda for this session

- 1 Unity AI Gateway
- 2 Unity AI Gateway Deep dive
- 3 Omnigent
- 4 Q&A

Unity AI Gateway: How it works?

The Unity AI Gateway architecture centralizes governance, routing, guardrails, observability, and cost controls for all AI traffic through a single, secure control plane.

01 AI assets are registered in Unity Catalog as securable services, so permissions and policies are managed like standard data assets.

02 Client requests hit Unity AI Gateway, which acts as the central control plane and secure front door for all AI interactions.

03 The gateway applies policy & traffic management, enforcing real-time access controls, guardrails, rate limits, budgets, and fallback routing.


04 Requests are forwarded to services across Databricks-hosted models, external providers, or MCP tools.

05 Usage, cost, and telemetry logs are written directly to Unity Catalog tables for simplified monitoring, audit, and debugging.




Unity AI Gateway: How it works?

Every request takes the same governed path



Democratize access to any AI model



Govern



Govern

Assets (who can use what) and interactions (what happens on each call).



Monitor

See every single call, tracking and understanding cost and decisions instantly.



Optimize

Control your enterprise spend while keeping AI integrations reliable.



Extend

Easily plug in your security, identity systems, and partner software stack.




AI assets are Unity Catalog securables

Same governance you built for data now covers models, MCPs, and agents.

Register once, govern it centrally: Each AI asset is a UC "service" securable (catalog → schema → service): Models, MCP, Skills, Agents

RBAC: grant EXECUTE (use), MANAGE (administer), CREATE SERVICE / CREATE CONNECTION (register)

Unified Governance: One governance model that extends to every next asset type. Same discovery, lineage, and audit as data.



Govern every **model** from one place

Every model, every provider: OpenAI, Anthropic, Bedrock, Foundry, Databricks-hosted and more; wrap an existing gateway too.

Frontier models, built in: Access leading models with no provider setup or capacity to manage.

Invoke with any API: Responses API, Messages API, Gemini API, provider passthrough, and batch, all through one governed endpoint.

Control access, rate limits, and credentials: govern per model for any principal; register keys once, encrypted and central, so callers use the model by name and never touch secrets.

The screenshot displays the Unity AI Gateway interface, which is used for governing and accessing AI models. The dashboard includes a navigation bar with tabs for Models, Beta, Providers, MCPs, Skills, and Agents. A prominent message states: "New! Model services now live in Unity Catalog with permissions, lineage, and cross-workspace governance built in." Below this, there are four key performance indicator (KPI) cards: "Agent Integrations" (Manage usage and access for external agents), "Featured models" (Frontier models hosted by Databricks), "Fastest response" (Output tokens per second), and "Highest Traffic" (Total input and output tokens in the last 7 days). The "Fastest response" and "Highest Traffic" cards list the top models. At the bottom, a table provides a detailed overview of the models.

Name	Configuration	Total tokens (7d)	Location	Created by	Last modified
claude-fable-5		71	system.ai	Databricks	Jun 8, 2026, 05:00 PM
claude-haiku-4-5		302.2K	system.ai	Databricks	Dec 18, 2025, 04:00 PM
claude-opus-4-1		0	system.ai	Databricks	Sep 3, 2025, 05:00 PM



Govern every **MCP and Tool** your agents call

Connect any tool: managed integrations for GDrive, Jira, Confluence, Slack, GitHub, SharePoint, plus Databricks-native (Genie, SQL) and your own custom servers, all in one inventory.

Register once, use anywhere: govern and audit it in one place, then let coding agents, Genie, and external frameworks call it across any workspace.

Fine-grained access, per tool: grant per server and enable/disable individual tools within it, so agents get only the tools they should.

Agents act as the user (OBO): tool calls can run with the requesting user's exact permissions via managed OAuth.

The screenshot displays the Unity AI Gateway dashboard. At the top, it says "Unity AI Gateway" with a subtitle "Govern and access your AI." and a "Learn more" link. A "Govern" dropdown menu is in the top right. Below the header, there are tabs for "Models", "Beta", "Providers", "Beta", "MCPs" (which is selected), "Skills", and "Agents".

The "Featured MCPs" section shows four cards: "Genie" (Ask natural-language questions against your data), "Atlassian" (Search Jira issues and Confluence pages), "GitHub" (Search code, manage issues and pull requests), and "Slack" (Send and search messages across channels).

Below this is a search bar "Search Servers" and a "Type: All" dropdown. A "+ Create MCP" button is on the right.

The main section is a table of managed MCPs:

Name	Description	Governance	Status	Location	Created By
atlassian			Active	system.ai	2a34bc1b-8708-4...
slack			Active	system.ai	2a34bc1b-8708-4...
js_svc_ext	example.com httpGe...		Active	main.tibicrai_ai	tibi.craiu@databric
xiaoyi_github_mcp			Active	main.default	xiaoyi.zhang@data
mcp			Active	main.max-roze	max.rozen@databr
gmail			Active	system.ai	2a34bc1b-8708-4...
google_drive			Active	system.ai	2a34bc1b-8708-4...



Share and discover **Skills** via central catalog

Publish once, reuse everywhere

Register a reusable capability so teams build on approved skills instead of rebuilding them.

Discover across the org

Central, searchable inventory of what is available.

Same controls

Access and audit as every other asset

The screenshot displays the 'Skills' tab in the Databricks central catalog. It features a sidebar with navigation icons, a top navigation bar with tabs for Overview, Models, Providers, MCPs, Skills (selected), and Agents. The main content area shows 'Featured skills' as a grid of six skill cards: SQL Analyst (128K invocations), Document Summarizer (96K), Data Visualizer (74K), Web Search (62K), Translator (45K), and Code Explainer (31K). Below the grid are filters for 'Search skills', 'Curated by: ALL', 'Use case: ALL', 'Status: ALL', and an 'Add skill' button. At the bottom is a table with columns: Name, Description, Use case, Invocations (7d), Users, Owner, and Status.

Name	Description	Use case	Invocations (7d)	Users	Owner	Status
SQL Analyst						Active
Document Summarizer						Active



Manage all **Agents** in one place

One catalog of every agent

Databricks-built and external (Salesforce Agentforce, Azure AI Foundry, and others), registered and discoverable in one place.

Capture usage and observability

Visibility into every agent activity: usage, traces, and metrics

Give each agent an identity

So its actions carry lineage and policies can be enforced on what it does

The screenshot shows the 'Agents' tab in the Databricks management console. It features a sidebar with navigation icons, a top navigation bar with tabs for Overview, Models, Providers, MCPs, Skills, and Agents. The main content area displays 'Featured agents' as a grid of cards, each representing an agent with its name, token usage (e.g., 92.4M), and a time period (7d). Below the grid is a search bar and three filter dropdowns for 'Curated by', 'Use case', and 'Status'. At the bottom, a table lists agents with columns for Name, Description, Use case, Users, Tokens (7d), and Owner.

Name	Description	Use case	Users	Tokens (7d)	Owner
github-pr-rev...				92.4M	
slack-channe...				68.2M	
salesforce-ac...				41.8M	
meta-ads-op...				28.9M	
marketo-cam...				18.3M	
data-analyst...				14.2M	



Govern at scale with Policy

Identity-aware content protection and access permissions

Content protection

Column masking

Mask sensitive columns

Mask any column tagged 'sensitive' to '*****' if agent is acting OBO user

Row filters

Filter rows for use cases

When agent acting OBO only return data for region = "US"

Service policies Beta

Filter interactions over MCPs, model services, ...

ASK before pushing code to github if agent is acting OBO a user

Access

GRANT policies Beta

Grant permissions using attributes

Grant execute to ml-team on all models tagged with "approved"

DENY policies PrP

Deny permissions using attributes

Deny `MANAGE ACCESS CONTROL` when agent is acting OBO a user



Manage permissioning with **Access Policies**

Write one rule, cover the whole class

Give ``analysts`` access to every model tagged ``lifecycle = production``, or ``data_scientists`` to all Anthropic models in ``system.ai``.

New assets are covered automatically

Anything that matches the tags is governed the moment it appears; no grant to add.

Govern by what a model is

Use custom tags or built-in ones like ``ai.model_creator`` and ``ai.model_family`` instead of listing objects one by one.

One framework across data and AI

The same ABAC that powers row filters and column masks on your tables.



```
CREATE POLICY
grant_anthropic_models
ON SCHEMA system.ai
TO `data_scientists`
EXCEPT `contractors`
GRANT EXECUTE FOR MODELS
WHEN has_tag_value(
  'ai.model_creator',
  'anthropic'
);
```



Govern agent interaction with **Service Policies**

Enforced in real time, on every call

Evaluated before the model runs (ON CALL) and after it responds (ON RESULT), across models, providers, and MCP tools.

Control flows and agent capabilities

Block PII and unsafe content, stop jailbreaks, restrict tool calls, and require approval for risky actions.

Set it once, apply everywhere

Write one policy and bind it across many services instead of configuring guardrails endpoint by endpoint.



Every request intercepted before execution. Every response inspected after.

ALLOW proceeds. **DENY** blocks with a reason. **ASK** pauses for human approval.



Types of Service Policies

Available Now

Built-in

Out-of-the-box, no authoring: block PII, unsafe content, jailbreak/prompt injection (request); catch hallucinations (response).

Custom, LLM-as-judge

A prompt-based classifier for nuanced content, e.g. deny responses that give financial advice.

Custom, scope by caller and tool

A SQL rule, e.g. DENY when a contractor's agent calls the GitHub "push" tool; ALLOW otherwise.

Contextual

Decision depends on what the agent touched: sensitive PII + "publish to website" = ASK; same on public data = ALLOW.

Coming Soon

Custom code

Arbitrary logic as a policy.

External / partner

Call a partner guardrail (Palo Alto, Zscaler, ...) and enforce its verdict inside the same policy.



Monitor



Govern

Assets (who can use what) and interactions (what happens on each call).



Monitor

See every single call, tracking and understanding cost and decisions instantly.



Optimize

Control your enterprise spend while keeping AI integrations reliable.



Extend

Easily plug in your security, identity systems, and partner software stack.



Usage and cost tracking

One cost view across every model, coding agent, and agent

One bill for all your AI


Token usage and spend across Databricks-hosted models, external providers, coding agents (Claude Code, Cursor, Codex, Gemini CLI), and custom agents.

Attribute spend to the source

Per-request, per-user, per-team, per-endpoint, and per-tag tracking in UC system tables to answer "what did this team spend on this model?"

Out-of-the-box Cost Analysis

Slice and analyze spend by coding agent, user, model, and tag instantly via pre-built dashboards.



Unified Trace Table

OpenTelemetry Table, for every request and response across all services

Complete

Models, MCPs, and agents all write to one table.
Multi-step sessions stay in one record instead of scattered tables.

OTel Compatible Traces

Each row is an OTEL span capturing trace, enabling precise agent debugging and security investigation.

Enforceable

Configure once. Default-on logging across all services with zero blind spots.

Open Standard

Built on OpenTelemetry, allowing Lakewatch, Genie, MLflow, and external SIEMs to consume data directly.

The screenshot shows the Databricks Governance console. The main panel displays the 'Traces' tab under 'Governance'. It includes a search bar, filters, and a table of traces. The table has columns for Time, Trace ID, Service, and Principal. The sidebar on the right shows a 'Recent Error' dialog with a query and a table of error occurrences.

Time	Trace ID	Service	Principal
06/12/2026, 11:29:59	8573e584f14042d96827d9ac7890db	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:29:41	6b2b8349f6d7127b80f0f9c17ac77b	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:29:28	417b75c3064667b2d489e575e84343	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:26:39	a3616494fc0487da30aa49e8b6af794	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:26:38	fc060e7455a4aaa6f626894191b1e2	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:26:37	3dfe3a286a048f8b9f0d4f83f3d20	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:26:30	7e6fb22d7544c52aca3a35769f741b6	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:26:29	b792f3b4f6734ee98cf3c61f96cc21	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:26:28	bab0b52e48b1445399c99508a9481848	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:26:18	c2889c70977a1aab8b80c596f3c85d4	LLM system.ai.claude-opus-4-8	aarushi.shah@databricks.com
06/12/2026, 11:19:52	8571725a56c04c93a62a38d7303b6f4c	databricks-claude-opus-4-8	arjun.dcurha@databricks.com
06/12/2026, 11:19:49	1b7464fecc6439a8e07c0ef84e41ffe	LLM databricks-claude-sonnet-4-8	arjun.dcurha@databricks.com

The sidebar on the right shows a 'Recent Error' dialog with a query and a table of error occurrences.

Recent Error ...

unity_gateway_otel_spans

How many errors did I get in the last few days, and what caused them?

5 steps

The table has a status struct with code and message fields. Let me query for errors in the last few days.

2 steps

You had 1,213 errors over the last two days (1,105 on Jun 11, 108 so far today). The top causes:

Error	Occurrences	Traces
rate_limit_exceeded	562	281
invalid_request	481	241
Access denied (policy)	55	55
internal_error	33	33
not_found	24	12
unauthenticated	8	4
Unknown policy handler system.ai.invoke_llm_judge	8	8
Agent cannot merge Pts	6	6

The vast majority (~86%) are rate limiting and invalid requests.

@ for objects, / for commands, + for history

+ @ Agent

Always review the accuracy of responses.



Your AI data, in your existing stack

All gateway's observability data lands as governed tables in Unity Catalog, so the tools you already run work on it directly.

SQL, dashboards, and BI

It is just governed Delta, so query it, chart it, and alert on it with the tools your teams already use.

Lineage


It sits in the same catalog as your data, so you get end-to-end data → model → agent lineage; swap a model and see downstream dependencies first.

Genie

Ask in natural language ("show all rate-limit errors on this endpoint in the last 24h"); it writes the SQL automatically.

Lakewatch and SIEMs

Point threat detection at the data (PII, prompt-injection, anomalies) and export cases to Jira or Tines



Optimize



Govern

Assets (who can use what) and interactions (what happens on each call).



Monitor

See every single call, tracking and understanding cost and decisions instantly.



Optimize

Control your enterprise spend while keeping AI integrations reliable.



Extend

Easily plug in your security, identity systems, and partner software stack.



Multi-AI cost management

Stop tokenmaxxing, start valuemaxxing

Budget any way you slice spend


By team, user, group, endpoint, or use case (via tags), across Databricks-hosted and external models alike.

Shared and per-user limits

A pooled team budget plus a per-user threshold each person gets, with overrides to lift caps for power users.

Act at the threshold

Alert, then hard-cap: block traffic within minutes of a budget being exceeded, with different actions (e.g. warn at 80%, block at 100%).



Create your Unity AI Gateway Budget

Governance

- Data
- Cost

Admin

- Workspaces
- Catalog
- Usage**
- User management
- Security
- Cloud resources
- Previews
- Settings

Usage

Consumption Cost Overview Preview **Budgets** Preview Consumption (Legacy)

Filters Filters com

Create budget

Name	Workspaces	Resource type	Tags	Shared threshold status	Per user threshold status
No budgets matched your search Try adjusting your filters					

Send feedback



Create your Unity AI Gateway Budget

Workspaces

Catalog

Usage

User management

Security

Cloud resources

Previews

Settings

Create budget

Scope

Name

conding-agent-budget

Workspaces

Select workspaces to include in this budget

test-cesar-workspace-1 X

Resource types

Cannot be changed after creating budget

Unity AI Gateway

Resource tags (optional)

Select tags to restrict budgets to specific types of usage. Usage with all keys applied will be tracked. If a tag key has multiple values, for each key at least one of the selected values applied will be tracked.

Select tags

Shared thresholds

Apply across all scoped spend

Monthly threshold

\$1,500.00

+ Add threshold

Per-user thresholds

Apply for each user, for their scoped spend. This threshold will be applied for all users.

Monthly threshold

+ Add threshold

ACME Account

Governance

Data

Cost

Admin

Workspaces

Catalog

Usage

User management

Security

Cloud resources

Previews

Settings

Usage > Budgets >

Coding Agent Budget

Spend is in list price, pre-credits, incl. add-ons

Edit budget

Overview

Shared spend

Per user spend

Filter by user

All statuses

Threshold

\$200.00

Monthly

Users

5 of 81 exhausted

User	Spend	% of threshold	Status
Wren Nakamura	\$895.99	448%	Exhausted
Emery Quintero	\$404.56	202%	Exhausted
Oakley Kovacs	\$296.25	148%	Exhausted
Dakota Diaz	\$207.81	104%	Exhausted
Clay Garcia	\$206.39	103%	Exhausted
Blair Chen	\$128.22	64%	Under budget
Jordan Hayes	\$123.93	62%	Under budget

Name	Spend	% of threshold	Status
Wren Nakamura	\$895.99	448%	Exhausted
Emery Quintero	\$404.56	202%	Exhausted
Oakley Kovacs	\$296.25	148%	Exhausted
Dakota Diaz	\$207.81	104%	Exhausted
Clay Garcia	\$206.39	103%	Exhausted
Blair Chen	\$128.22	64%	Under budget
Jordan Hayes	\$123.93	62%	Under budget

Monitor and Control Budgets Through Unity AI Gateway

- Set alerts and hard caps for workspaces and users
- Overrides for users and groups (e.g., power users)
- Cost observability dashboards
- System table integration

Account name

Account Console

Governance

Cost

Admin

Workspaces

Catalog

User management

Usage

Cloud resources

Previews

Settings

Cost Governance > Budgets >

Account

Create budget

Scope

Name

claude-code-gateway

Workspaces

Cannot be changed after creating budget

development

Resource types

Cannot be changed after creating budget

Unity AI Gateway

Resource tags (optional)

Select tags

☒ Include external model usage

Shared thresholds

Shared by all users

Users	Monthly threshold	When exhausted
(All users)	\$5,000.00	☒ Alert (1)
(All users)	\$10,000.00	☒ Block ☒ Alert (2)

+ Add threshold

Per-user thresholds

Applied to all users, tracked separately per user

Users	Monthly threshold	When exhausted
(Default users)	\$200.00	☒ Block
engineering	\$300.00	Inherited from default
power-users	\$500.00	Inherited from default
Khai Nguyen	\$600.00	Inherited from default

+ Add override

Permissions

Cannot be changed after creating budget

Account admins	Can manage
Workspace admins	None

Cancel

Create

Governance

Data

Cost

Admin

Workspaces

Catalog

User management

Usage

Cloud resources

Previews

Settings

Cost Governance > Budgets >

Account

claude-code-gateway

Spend is in list price, pre-credits, including add-ons


Edit budget

Overview Shared spend Per-user spend

Filters

Search users

Showing spend this month. Users will display as they incur usage.

Exhausted users
3Monthly users
232

Name	Status	Spend	Per-user threshold	User group	% of threshold
Khai Nguyen	Exhausted	\$535.00	\$500.00	power-users	109%
Sadhana Bala	Exhausted	\$500.00	\$500.00	power-users	100%
Eliza Stone	Exhausted	\$500.00	\$500.00	power-users	100%
Ricardo Baca	Under budget	\$286.45	\$500.00	power-users	84%
Kayla Singh	Under budget	\$279.94	\$500.00	power-users	78%
Armando Smith	Under budget	\$268.43	\$300.00	engineering	64%
Tristan Leal	Under budget	\$261.50	\$300.00	engineering	56%
Lara Croft	Under budget	\$250.75	\$300.00	engineering	55%
Jamal Adebayo	Under budget	\$210.00	\$300.00	engineering	44%
Priya Desai	Under budget	\$200.00	\$300.00	engineering	19%



Budget “eng-budget” threshold exhausted for user “Khai Nguyen”

Spend of user “Khai Nguyen” in user group “power-users” has exhausted their per-user threshold on the budget “eng-budget”.

[Review in Databricks](#)

Budget name	eng-budget
User	Khai Nguyen
User group	power-users (Override)
Current spend	\$10,005.56
Per-user threshold	\$10,000.00 (Override)
Budget period	May 1, 2026 – May 31, 2026
Percent of period elapsed	9.6%

Budget scope

Workspaces	production, development, staging, prod-analytics, testing, example, long-example, workspace-1, workspace-2, workspace-3, workspace-3
Resource types	Unity AI Gateway
Tags	-



You’re receiving this because a Databricks admin added you as a recipient on this budget. To stop receiving these notifications, edit the budget or contact your admin.

Budget ID: 1bca0da2-7097-43ea-a862-350500bcc7be

Smart Routing

Optimize cost without cutting access

- Simple tasks → faster, lower-cost models
- Complex tasks → more capable models
- Automatically falls back if a model is unavailable
- Reduces cost without sacrificing quality



Traffic management and reliability

Keep AI reliable and within limits

Rate limits


Cap tokens/queries per minute per user, endpoint, or key, so no one workload starves the rest.

Fallback

If the primary model degrades or goes down, traffic auto-routes to a backup. No code changes.

Load balancing / traffic splitting

Spread load across providers and endpoints for resilience and scale.



Extend



Govern

Assets (who can use what) and interactions (what happens on each call).



Monitor

See every single call, tracking and understanding cost and decisions instantly.



Optimize

Control your enterprise spend while keeping AI integrations reliable.



Extend

Easily plug in your security, identity systems, and partner software stack.



Build agents on AI Gateway with Omnigent

Meta harness for agent-first developers

Composition.


Combine multiple harnesses, like Claude Code and Codex, without rewriting code.

Control.

Stateful, contextual policies that track agent actions and enforce guardrails.

Collaboration.

Pair program or code review with colleagues and agents in the same session.



Open Partner Ecosystem for AI Governance

- New integrations across AI security, identity, observability, and agent governance
- Open partners integrated directly into Unity AI Gateway
- 14+ partners and growing

Unity AI Gateway Partner Ecosystem

ALICE

CROWDSTRIKE



HIDDENLAYER



NETSKOPE

NOMA



OBSIDIAN



OKTA



OPENLAYER

PALOALTO NETWORKS

PINGIDENTITY



SAVIYNT



ZSCALER

CYERA



SAILPOINT

+ many more partners



Need on-prem or deep customization?

Use MLflow AI Gateway

When you need to self-host


Strict data-residency or on-prem mandates: run the open-source MLflow AI Gateway in your own environment.

When you need to customize

Open source architecture allows you to extend routing, providers, and logic beyond what the managed gateway exposes.

Same shape, carries over

Unified provider interface, routing, and OpenAI-compatible APIs ensure what you build maps back seamlessly.



Deployment and Migration



AI Gateway landscape

This is the July launch and what we'll take to GA.

	Legacy AI Gateway	Unity AI Gateway (Workspace Models)	Unity AI Gateway (Unity Catalog Services)
Internal Versions	v1	v2	v3
Scope	Workspace	Workspace	Unity Catalog
Location in product	Model Serving tab	Separate left nav	Separate left nav
Governs	Models only	Models & MCPs	Models, MCP, Agents, Skills
Status	GA	Beta, being retired	Beta (launched first week of July), moving to GA soon.
Enrollment	Available in all workspaces where Model Serving is present	Beta rollout: default on for AWS Ring 3, self-serve for AWS Ring 2, manual enrollment for all others	Automatically rolled out to all accounts that are enrolled to Unity AI Gateway (Workspace Models) Beta.



Start Everything New on Unity AI Gateway.

One simple rule: all new projects and activity begin on Unity AI Gateway.

Existing endpoints won't break

Anything running today keeps working.

Unity AI Gateway(workspace, v2) is deprecating soon

Move to Unity AI Gateway (V3).

Legacy AI Gateway continues for now

Deprecation details will be shared later.

Unity AI Gateway(v3) GAing early Q3

The version to build on going forward.

The screenshot displays the Unity AI Gateway dashboard. On the left is a sidebar with navigation options: Analytics (Dashboards, Genie Agents, Alerts, Visual Data Prep), Lakehouse (SQL Editor, Queries, Query History, SQL Warehouses), Data Engineering (Runs, Data Ingestion), AI/ML (Playground, Agents, Experiments, Features, Models), and AI Gateway. The main panel is titled 'Unity AI Gateway' with the subtitle 'Govern and access your AI. Learn more'. Below this are tabs for Models, Beta, Providers, MCPs, Skills, and Agents. A notification states: 'New! Model services now live in Unity Catalog with permissions, lineage, and cross-workspace governance built in. Learn more'. Three cards are visible: 'Agent Integrations' (manage usage and access for external agents), 'Featured models' (Frontier models hosted by Databricks), and 'Fastest response' (Output tokens per second). Below these is a search bar and filters: 'Owned by me', 'Created by', 'Configuration', and 'Location'. A table lists available models:

Name	Configuration	Total tokens (7d)	Location	Created by
claude-fable-5		71	system.ai	Databricks
claude-haiku-4-5		301.6K	system.ai	Databricks
claude-opus-4-1		0	svstem.ai	Databricks



Migrating to Unity AI Gateway

Unity AI Gateway (UC)

GA Early Next Quarter

The ultimate target gateway version everyone moves to.

Unity AI Gateway (Workspace)

Retires Fast

New creation already blocked.

Existing endpoints wind down shortly after GA (targeting ~1 month).

Legacy AI Gateway

Gentle Wind-Down

Longer runway since it is a product

Targeting ~6 months from GA to end of support. Non-LLM serving endpoints are unaffected.

Assisted Migration

Guided Transition

Will provide guided in-product path + resources + maybe traffic rerouting.

Most routing and governance settings carry over seamlessly.



Bring every developer in through your IdP

AI Gateway is per-user by identity — so first, get your people into Databricks.

Every gateway user is a Databricks user

Models, coding agents, and MCP servers resolve real identities. Users are provisioned by syncing from your IdP.


Sync via AIM (Preferred) or SCIM

AIM: Account-level JIT provisioning at login via SSO. Perfect for directory-wide "all users" access.

SCIM: Scheduled push sync. Use as fallback or when group-level control is required.

Identity does the work

Once provisioned, users authenticate via SSO/OAuth or PAT, enabling usage tracking, cost attribution, and rate limiting.



Secure the network path

The same Databricks networking controls you already use, applied to AI traffic.

Ingress

Front-end Private Link keeps user & API traffic off the public internet. IP access lists limit access to approved networks.

Egress

Serverless egress control applies a deny-by-default posture. Outbound Private Link routes traffic privately.


Encryption

Traffic is encrypted in transit (TLS 1.2+).

Control-plane-to-compute traffic stays on the cloud backbone.

Audit

Requests and responses flow through the gateway and are logged to Unity Catalog Delta tables for review.



AI Guardrails

Prevent data leakage and unsafe requests/responses.


The screenshot shows the Databricks AI Gateway interface for a model named "marketing-team-gpt-5-4". The interface includes a search bar at the top, a navigation menu with "Overview", "Guardrails", and "Metrics", and a "Delete endpoint" button. The "Overview" tab is selected, showing the model's status as "Active" and its URL. A "Chat in playground" button is also visible. The "Model" section displays the primary model "system.ai.databricks-gpt-5-4" with a "Pay-per-token" pricing model. Below this, there is an "Add another model" button and an "Add fallback" button. The "Gateway Endpoint Details" section on the right shows the endpoint was created by "amber.roberts@databricks.com" on May 13, 2026, at 04:55 PM EDT. It also includes options for "Usage tracking" (set to "system.ai_gateway.usage"), "Inference table" (with a "Set up" button), and "Rate limits" (with a "Set up" button). The "Tags" section is empty, and there is a link to "Applicable model terms". At the bottom, there is a "View starter code" button and a dropdown menu for "MLflow Chat Completions". A note at the bottom states: "# To get a DATABRICKS_TOKEN, click the 'Generate Access Token' button or follow https://docs.databricks.com/en/dev-tools/auth/pat.html".



One view across all AI activity

A single pane of glass for everything your AI does in production.

- **Every interaction logged:** Requests, responses, tokens, latency, errors, user identity, and guardrail outcomes, all in a UC-governed Delta table.
- **Query with tools you already use:** SQL, dashboards, jobs, alerts. Join AI logs with business data for richer insights.
- **Access governed by Unity Catalog:** Same permissions as any other table, with fine-grained control over who sees what.



AI Gateway inference tables (auditing and quality)

Monitor quality and audit for compliance

Payload logging
(inference tables) = logs request and response to UC delta table

Share with developers to monitor quality and use to audit for compliance.

The screenshot shows the Databricks AI Gateway console for a gateway named 'archika-gateway-cuj'. The 'Gateway' section is highlighted with a red box, showing the following configuration:

- Use monitoring: [system.serving_endpoint_usage](#)
- Dimension table: [system.serving.served_entities](#)
- Inference tables: [archika-catalog.default.archika-gateway-cuj_payload](#)

Other details visible include: Endpoint state: Ready, Created by: archika, Last modified: 11 days ago, URL: (empty), Task: Chat, Tags: Add, Rate limits: Not enabled, Input guardrails: PII detection, Safety (3 valid topics, 2 invalid keywords), Output guardrails: PII detection, Safety (3 valid topics, 2 invalid keywords). The 'Active configuration' table below shows the gateway is using the 'gpt-4o-mini' model from 'OpenAI' with an 'openai_api_type' of 'openai'.



Payload Logging (aka "Inference Tables")

All incoming requests and outgoing responses are stored in a `delta` table in the user's account. For now, the user will have to unpack the JSON.

NOTE: This is the only table that contains the request and response logging.

Raw Payloads

request_date	databricks_request_id	request_time	status_code	execution_duration_ms	request	response	served_entity_id	requester
2024-08-22	040f49e-c208-447a-8...	2024-08-22 16:35:...	200	4209	{ "messages": [{"role": "user", "content": "Can you make ..."}]}	{ "id": "chatcmpl-9254...", "text": "..." }	9bf59aadb23b47b68...	archiki
2024-08-22	7381a640-342e-4383-8...	2024-08-22 16:33:...	200	1905	{ "messages": [{"role": "user", "content": "How old is the ..."}]}	{ "id": "chatcmpl-9252...", "text": "..." }	9bf59aadb23b47b68...	archiki
2024-08-22	a6e3204d-d8dc-4571-b...	2024-08-22 16:35:...	200	6358	{ "messages": [{"role": "user", "content": "How to change ..."}]}	{ "id": "chatcmpl-9254...", "text": "..." }	9bf59aadb23b47b68...	archiki

Unpacked Payloads

databricks_request_id	request_time	status_code	requester	messages	choices
2e1df2e6-30fa-4909-b21c-504f22738925	2024-08-21 19:48:16.839	200	kasey	[{"role": "user", "content": "..."}]	[{"text": "..."}]
caf4d1c5-8db9-43fd-96d0-ed6fc60b749d	2024-08-21 19:50:03.082	200	kasey	[{"role": "user", "content": "..."}]	[{"text": "..."}]



AI Gateway inference tables (auditing and quality)

Catalog Explorer > stable_classic_6kvr7_catalog > cesar_cordoba >

llm_testing_payload

Overview Sample Data Details Permissions Policies History Lineage Insights Quality

AI Suggested Description

The table tracks model serving requests and their details. It includes data such as request dates, client identifiers, execution analyzing request trends, and debugging issues related to request or response errors.

AI Gateway inference tables simplify monitoring and diagnostics for models by continuously logging serving request inputs and responses

Filter columns...					
Column	Type	request	response	served_entity_id	requester
databricks_request_id	string	{"messages":[{"content":"What is Databricks?","role":"user"},"temperature":0.0]}	> {"id":"chatcmpl_5c0f042-8838-4a22-8dfd-73ee330479c","object":"chat.completi...	47569d8154984ddcaaccfada5080d5f	cesar.cordoba@data
request_date	date		> {"id":"chatcmpl_16aff7a8-7405-4759-89a8-b73c13b116b5","object":"chat.completion.chunk","created":1769780086,"model":"b73c13b116b5","choices":[{"index":0,"message":{"role":"assistant","content":"Excelente pregunta! La lucha final entre Aladdin y Jafar es un momento emocionante y épico en la historia de Aladdin. Aquí te explico cómo logra Aladdin derrotar a Jafar y recuperar la piedra filosofal:\n\nLa lucha en el jardín secreto\n\nAladdin y Abu se encuentran en el jardín secreto, donde se encuentran con Jafar, que está dispuesto a sacrificar la piedra filosofal para obtener el poder de la inmortalidad y la riqueza. Aladdin se da cuenta de que Jafar está utilizando la piedra para crear un trampolín que le permita viajar a través del tiempo y el espacio, y que se enfrenta a la princesa Jasmine, que se encuentra atrapada en un compromiso con Jafar.\n\nLa batalla épica\n\nAladdin y Jafar se enfrentan en una lucha épica, con Aladdin utilizando sus habilidades y habilidades para defenderse y derrotar a Jafar. Aladdin utiliza su astucia y su ingenio para engañar a Jafar y distraerlo, mientras que Abu lucha contra los guardias que se han convertido en criaturas fantasmales.\n\nEl trampoline\n\nMientras tanto, Aladdin y Abu logran desbloquear el trampolín y llegar al corazón del jardín secreto. Allí, encuentran la piedra filosofal y se enfrentan a Jafar, que está dispuesto a sacrificar todo para obtener el poder. Aladdin se enfrenta a Jafar en una lucha épica, utilizando todas sus habilidades y habilidades para derrotar a Jafar.\n\nLa victoria final\n\nFinalmente, Aladdin logra derrotar a Jafar y recuperar la piedra filosofal. Jafar se sacrifica para evitar que la piedra se quede en manos de Aladdin, y Aladdin se convierte en el dueño de la piedra filosofal. La piedra filosofal luego se convierte en un símbolo de la libertad y la justicia, y Aladdin se convierte en un héroe para la ciudad de Agrabah.\n\nLa unión con la princesa Jasmine\n\nAl final, Aladdin y Jasmine se casan y se convierten en una pareja feliz. La historia de Aladdin se convierte en una leyenda que se transmite de generación en generación, inspirando a otros a buscar la libertad y la aventura.\n\nEspero que hayas disfrutado de esta explicación de la lucha final entre Aladdin y Jafar. ¿Times alguna pregunta o comentario sobre la	4973ba6814f940bebb1dc7f68dd3b7dd	cesar.cordoba@data
client_request_id	string	> {"messages":[{"role":"user","content":"cuentame la historia de aladin"},"role":"assistant","content":"Claro! Me encan...			
request_time	timestamp	{"messages":[{"role":"user","content":"cuentame la historia de aladin"},"max_tokens":5000,"stream":true]}	> {"id":"chatcmpl_416f430d-ba56-4b90-a978-8a0640144d5","object":"chat.compli...	4973ba6814f940bebb1dc7f68dd3b7dd	cesar.cordoba@data



Governance

Secure controls over who gets access and how much

- ❑ Apply **rate limits** per Endpoint or/and per user
- ❑ **Tokens** and **queries** per minute rate limits
- ❑ Set **user permissions** for served endpoint access

Set guardrails to prevent the model from interacting with certain types of content. [Learn more](#)

☒ **Rate limits**
Enforce request rate limits to manage traffic for this endpoint.

QPM ⓘ	
🔗 Endpoint ⓘ	No limit Reset
👤 User (Default) ⓘ	No limit Reset
👤 Archika Dogra (archika.dogra@databricks.com)	4 ×
Select user, group, or service princ... ▾	No limit Add

[+ New](#)

Serving endpoints >

archika-gateway-cuj

Endpoint state: Ready

Created by: 1

Last modified: 11 days ago

URL: [Permissions](#) [Edit](#) [Query](#)

Task: [Chat](#)


Tags: [Add](#)



Keep AI reliable when models fail

Keep agent quality consistent, even when models go down.

- **Fallback routing:** If the primary model degrades, traffic auto-routes to a backup. No code changes.
- **Load balancing:** Distribute traffic across providers for resilience and consistent response quality.



What's Coming: Coding Agents Gateway Updates


Help admins govern and roll out coding agents across the org | Q2 2026

Coding gateway rollout: Roll out multi agent coding agent across your org with OAuth, MCP integration, and centralized config management.

Per-user spend controls: Set budgets and spend limits per developer across all coding tools in UC-governed usage tables.

More OSS coding tools: Hermes, OpenCode, Pi.

Web search support for Claude: Enable Claude code to use web search



Track and control AI spend

Full visibility into AI spend with built-in budget controls.

- **Token-level cost attribution:** Per-request, per-user, per-endpoint tracking in UC-governed inference tables.
- **Budgets and spend limits:** Set thresholds per team, app, or endpoint. Alerts fire and access restricts at threshold.
- **Coding agents included:** Claude Code, Cursor, Gemini CLI all route through the same gateway. One cost view for every AI tool your teams use.

Total Policies

5

Exceeded

2

5 users blocked

With Enforcement

2

blocking or throttling

Q Search budgets...

NAME	WORKSPACES	TAGS	USERS & GROUPS	STATUS	WHEN EXCEEDED
default-monthly-budget	All	—	all-users per-user · monthly	4 of 11 exhausted	Block Access
power-user-budget	All	—	power-users per-user · monthly	1 of 4 exhausted	Block Access
prod-design-budget	All	—	product-design shared · weekly	15% spent	Warn Only
dev-gateway-budget	development	—	—	38%	Warn Only
chatbot-weekly-budget	All	chatbot	—	52%	Warn Only



Centralized usage tracking across all GenAI models (AI/BI dashboard)

Unity AI Gateway

Govern and access LLM and MCP endpoints. [Learn more](#)

LLMs **Beta** MCPs

Coding agents
Manage usage and access for coding agents

- Codex
- Cursor
- Gemini
- Other Integrations

Featured models
Frontier models hosted by Databricks

- GPT-5.4
- A Claude Opus 4.6
- Gemini 3.1 Pro
- GPT OSS 120B

Fastest response
Output tokens per second

1.	GPT OSS 120B	278.6
2.	GPT-5.4 Mini	94.8
3.	Claude Opus 4.7	84.6
4.	Claude Opus 4.5	51.4

Highest Traffic
Total input and output tokens in the last 7 days

1.	Claude Opus 4.5	186.8M
2.	Claude Opus 4.6	103.2M
3.	Claude Sonnet 4.6	53.9M
4.	Claude Opus 4.7	23.7M

☐ Owned by me

Created by: Databricks

View Dashboard

+ AI Gateway Endpoint

Gateway features

Name	Rate limits	Usage tracking	Inference table	Total tokens (7d)	Created by	Last modified
databricks-qwen35-122b				0	Databricks	Apr 22, 2026, 08:00 PM
databricks-claude-opus-				23.7M	Databricks	Apr 15, 2026, 08:00 PM
databricks-claude-opus-				186.8M	Databricks	Apr 14, 2026, 05:48 AM

SEE IT IN ACTION

Demo

Agenda for this session

- 1 Unity AI Gateway
- 2 Unity AI Gateway Deep dive
- 3 Omnigent
- 4 Q&A



&



Omnigent

Deploying Coding Agents at Enterprises is **Hard**

Security



Data privacy is a major concern - sensitive enterprise data may be used for training

Control



Admins are heavily involved in managing complex contracts and managing users/costs

Admin controls and observability are fragmented across different tools

Choice



The “best” coding tool keeps changing as models evolve

Committing to one coding tool is not the best option


What is a harness?



MODEL

*The AI brain
thinks, can't act.*


Ex: Claude 4.7,
GPT-5, Gemini



HARNESS

*Enables the brain
to take action.*

Ex: Claude Code,
Codex, Cursor



AGENT

*Understands
information and
take actions to
achieve a goal.*



We need a new, open layer above the
harness level.

The meta-harness



Omnigent

Meta harness for agent-first developers

Secure.


Run coding agents at full speed. Your policies decide when to pause, not you.

Portable.

Centrally distribute your MCPs, memory, skills, and artifacts across harnesses or devices (CLI, mobile, etc)

Collaborative.


Pair program or code review with colleagues and agents in the same session.




Managed on DB: TBC
Open-source: GA

Omnigent

Databricks' secure, portable,
collaborative **meta harness**.




Omnigent on Databricks



Competitive Landscape

Omnigent works with a customer's existing harnesses to maximize their experience



- The meta-harness layer that works with other harnesses
- No enterprise-ready, production-grade alternative



Let's Look a Little Closer:

ClaudeCode (Anthropic)

Omnigent competes with Anthropic / Claude Code by **sitting one layer above the coding agent**:

- Claude Code is the **strong first-party agent / assistant**.
- Omnigent is the **portable control layer** for running, supervising, sharing, and governing Claude Code alongside other agents.

Codex (OpenAI)

While developers use Codex, enterprises can use Omnigent as the **layer that runs Codex alongside** Claude Code and custom agents:

- OpenAI Codex specializes in long horizon, multi session **engineering** tasks.
- **Omnigent adds features that enterprises need:** shared sessions, policies, context, and collaboration.

OpenClaw, Pi, GasTown, etc


Omnigent is the **portable control layer** for running, supervising, sharing, and governing agents **across all services**:

- Pi is portable multi model terminal usage → **heavy setup required**
- Gas Town is local multi agent orchestration → **no managed hosting**
- OpenCode/OpenUI is multi-surface coding UX → **no governance**







Omnigent Meta-Harness



Omnigent Meta-Harness (OSS)



Hosted Omnigent Meta-Harness




Bringing various models under one roof

One Contract, Any Model, No Lock-in



What Omnigent gives you




Composition

- Combine agents using multiple harnesses / LLMs
- Switch agent harness



What Omnigent gives you




Composition

- Combine agents using multiple harnesses / LLMs
- Switch agent harness 

Collaboration

- Share and co-drive live sessions with other people
- Comment on files / docs 

What Omnigent gives you



Composition

- Combine agents using multiple harnesses / LLMs
- Switch agent harness



Collaboration

- Share and co-drive live sessions with other people
- Comment on files / docs



Control

- Powerful **contextual** policies for security and cost
- Strong OS sandbox



Omnigent

Meta harness for agent-first developers

Omnigent on Databricks is in Beta

[Public docs [AWS](#) | [Azure](#)]

Databricks provides a fully managed version of Omnigent, including:

- A Databricks-operated Omnigent server that integrates with your workspace's identity provider.
- Model access through the [Foundation Model APIs](#) and [AI Gateway](#).
- Databricks Sandboxes for secure, collaborative agent coding and knowledge work. Databricks Sandbox is available in [select regions on AWS](#).

In beta on Databricks, you will pay for models through Unity AI Gateway and any compute you use for Lakebox. Other charges have not yet been determined.

Competitive landscape

Major point: Omnigent is a new category of product. There are no enterprise-grade competitors at this time.

Other adjacent players:


- Coding agents (Claude Codex, Codex, Gemini CLI)
 - These are proprietary to whoever developed them. Omnigent is open.
 - You want more than one harness for the best quality and cost management.
- IDEs (Cursor, Windsurf)
 - These are still tied to the machine they run on.
 - No mobile or portable experience.
- Open source harnesses or tools (Pi, OpenUI)
 - DIY, must run and own yourself. What if they stop being supported?

Internal FAQ: https://docs.google.com/document/d/14gg_2j5TwyjfDjGATuEHHgC0JIETJ-CVF0H6Xe3a85E/edit?tab=t.0



Omnigent demo

UAIGW – Omni and MCP policies



Agenda for this session

- 1 Unity AI Gateway
- 2 Unity AI Gateway Deep dive
- 3 Omnigent
- 4 Q&A

Frequently asked questions



What is Unity AI Gateway?

Unity AI Gateway is Databricks' multi-AI management and runtime governance layer for enterprise AI systems, governing model calls, agents, MCP servers, AI tools, and autonomous workflows within Databricks or externally.



How is Unity AI Gateway different from Unity Catalog?

Unity Catalog governs the AI assets themselves: data, models, agents, functions, MCP servers, and policies. Unity AI Gateway governs runtime interactions, enforcing policies defined in Unity Catalog as well as guardrails, identity, and observability during every AI request and agent action.



Can Unity AI Gateway govern external models and AI systems?

Yes. Unity AI Gateway supports frontier and open-source models, external agents, MCP servers, and third-party frameworks with consistent governance controls across heterogeneous AI environments.



How does Unity AI Gateway help control AI cost and spend?

Unity AI Gateway provides centralized visibility and control over AI usage, including token monitoring, real-time rate limits, hard spend caps, budget policies, and spend attribution across teams, applications, agents, and model providers.



How does observability and monitoring work?

Unity AI Gateway captures traces, payload logs, audit events, and tool interactions across AI workflows, enabling monitoring, investigation, and SIEM integrations through MLflow Tracing and Lakewatch.



Q&A